

Lecture on Machine Learning



Maria J. Molina, Assistant Professor

University of Maryland, College Park

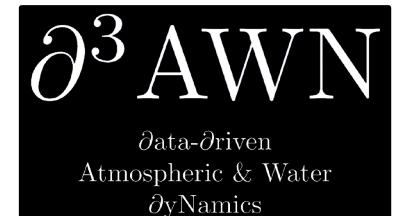
mjmolina@umd.edu; mariajmolina.github.io



Tom Beucler, Assistant Professor

Université de Lausanne

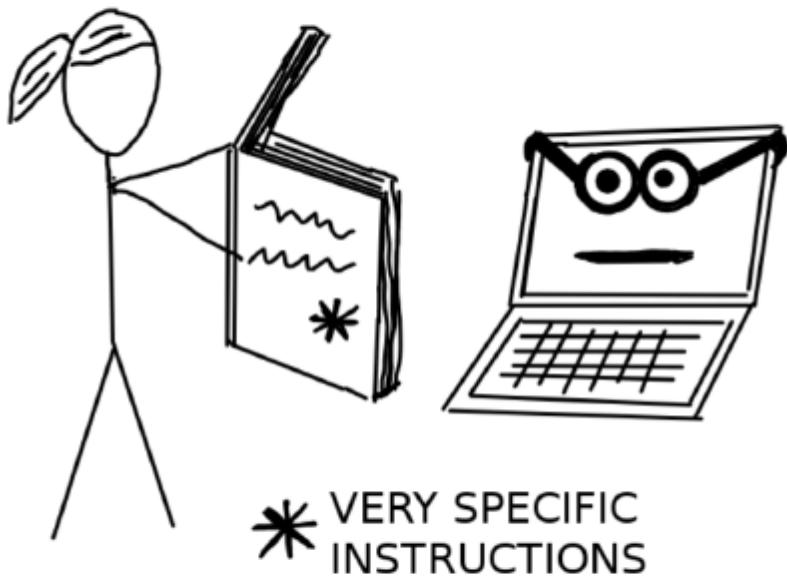
tom.beucler@unil.ch; wp.unil.ch/dawn/



Learning Outcomes (Profs. Molina and Beucler)

1. Frame an ML problem (e.g., regression or classification) for your specific application.
2. Choose appropriate data preprocessing steps.
3. Understand how ML models are trained.
4. Define physics-guided ML and apprehend the diversity of available approaches.

Without Machine Learning



With Machine Learning

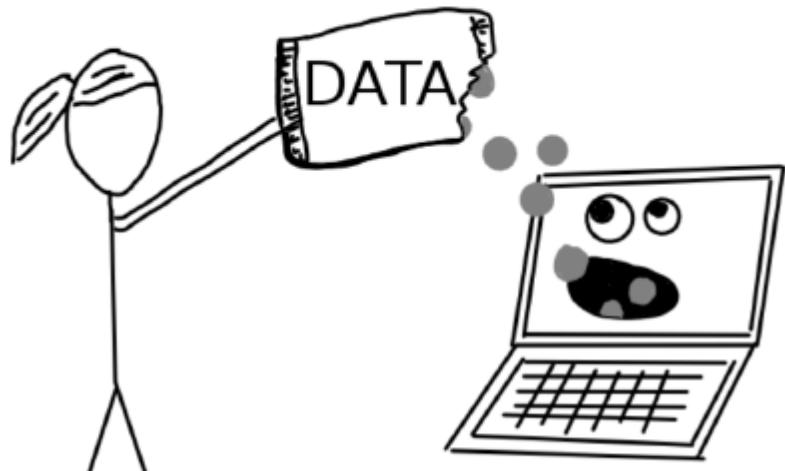
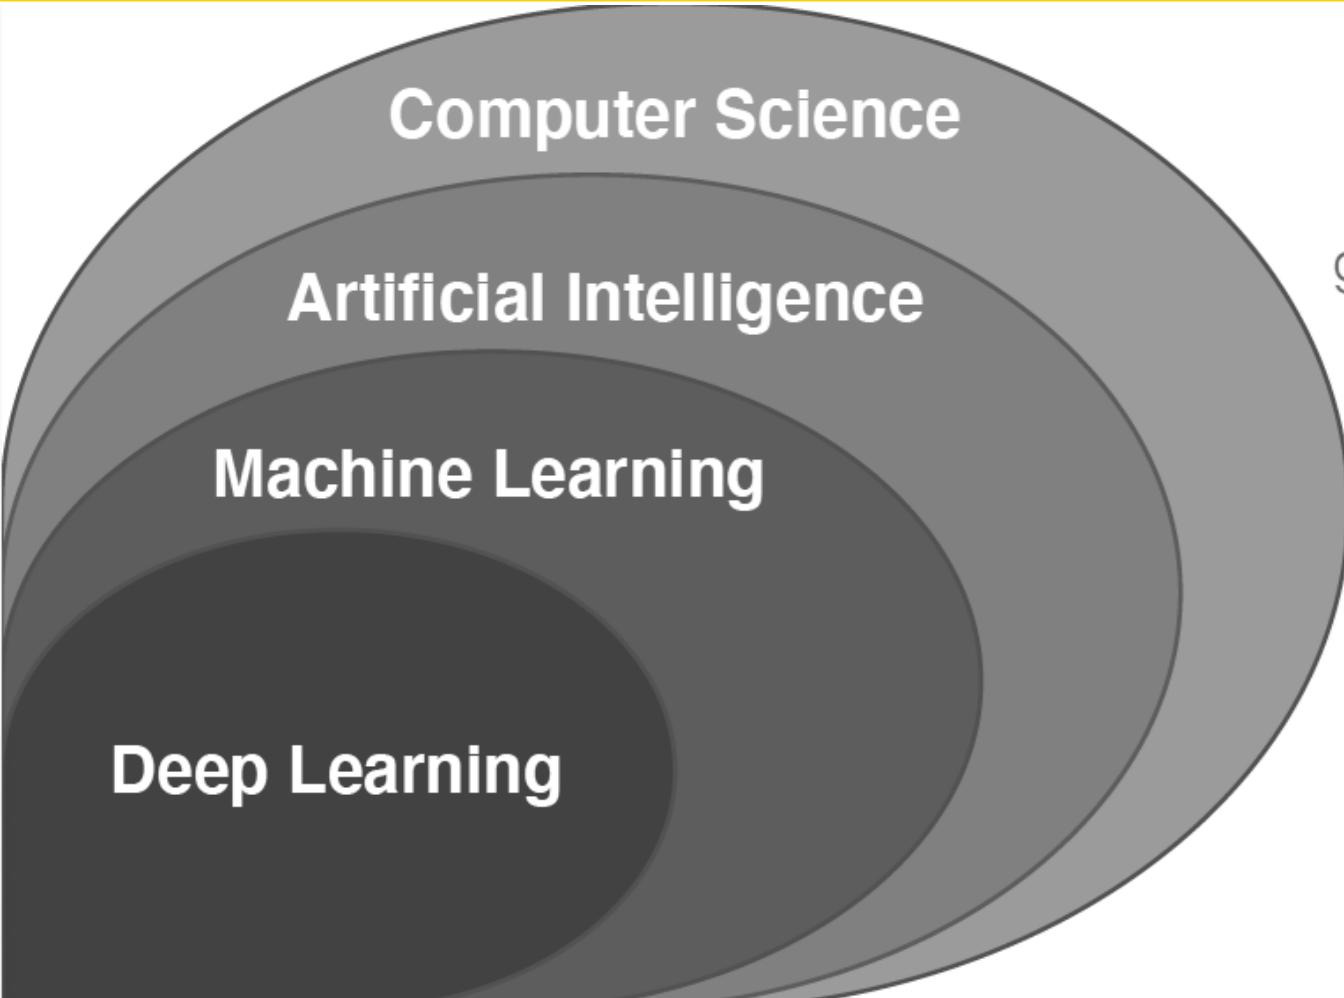
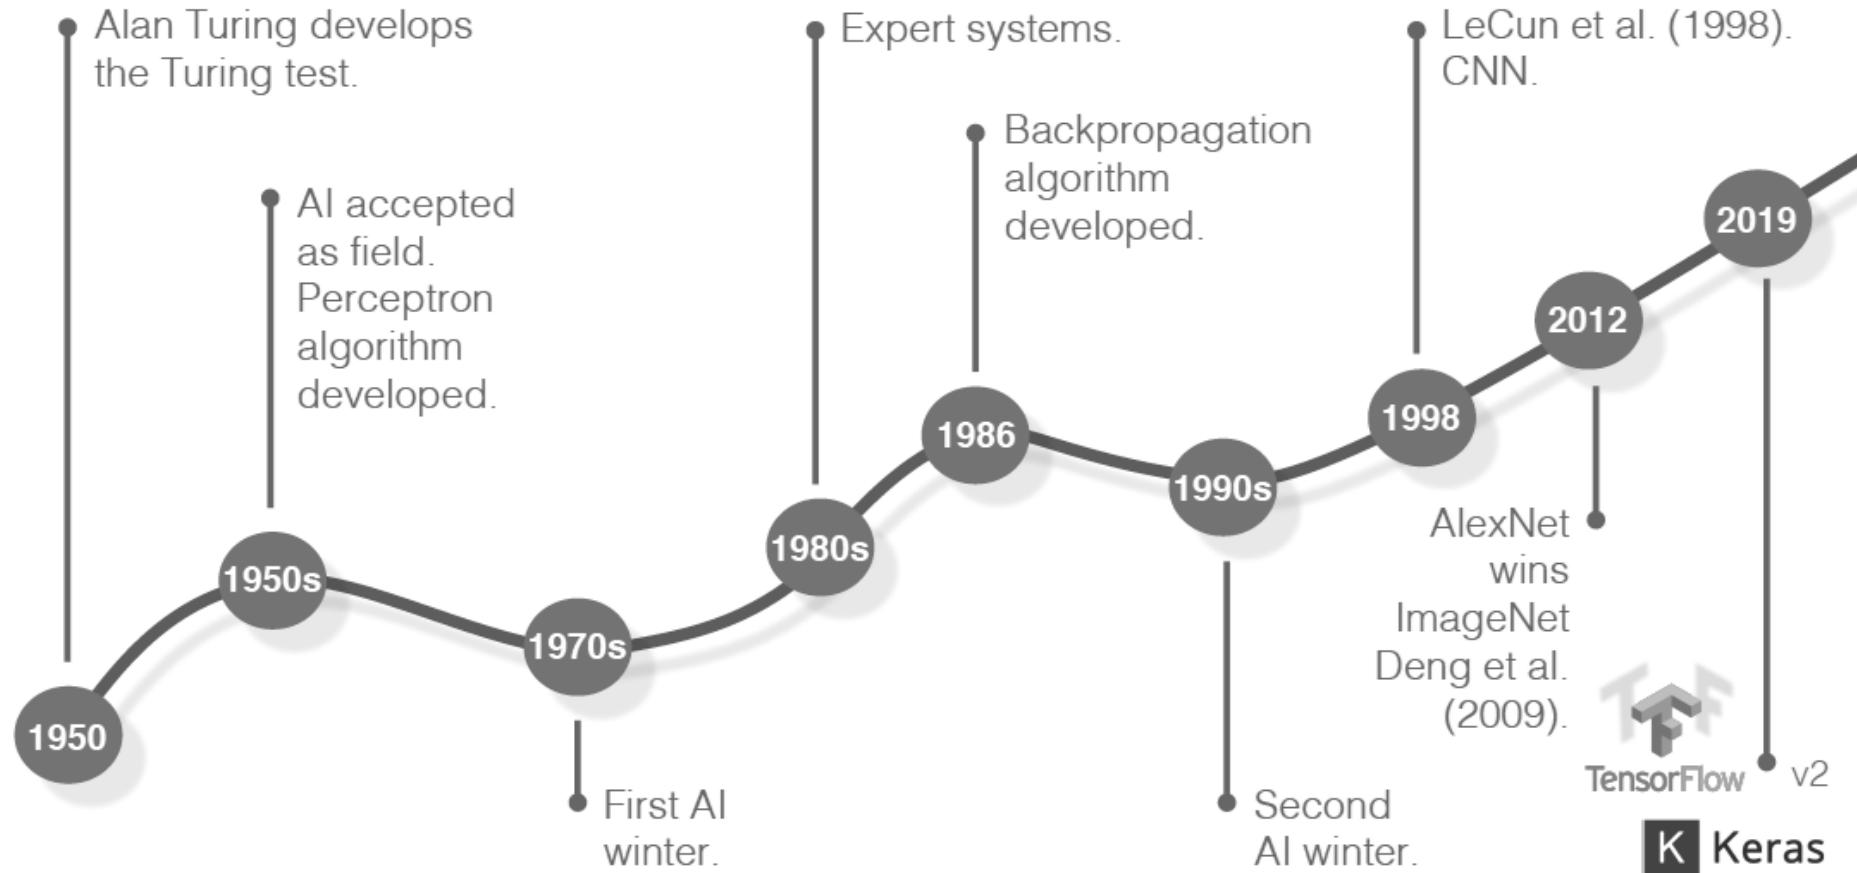


Figure from Interpretable Machine Learning (Molnar, 2019).



“Field of study that gives computers the ability to learn without being explicitly programmed.”

1959, Arthur Lee Samuel defined ML

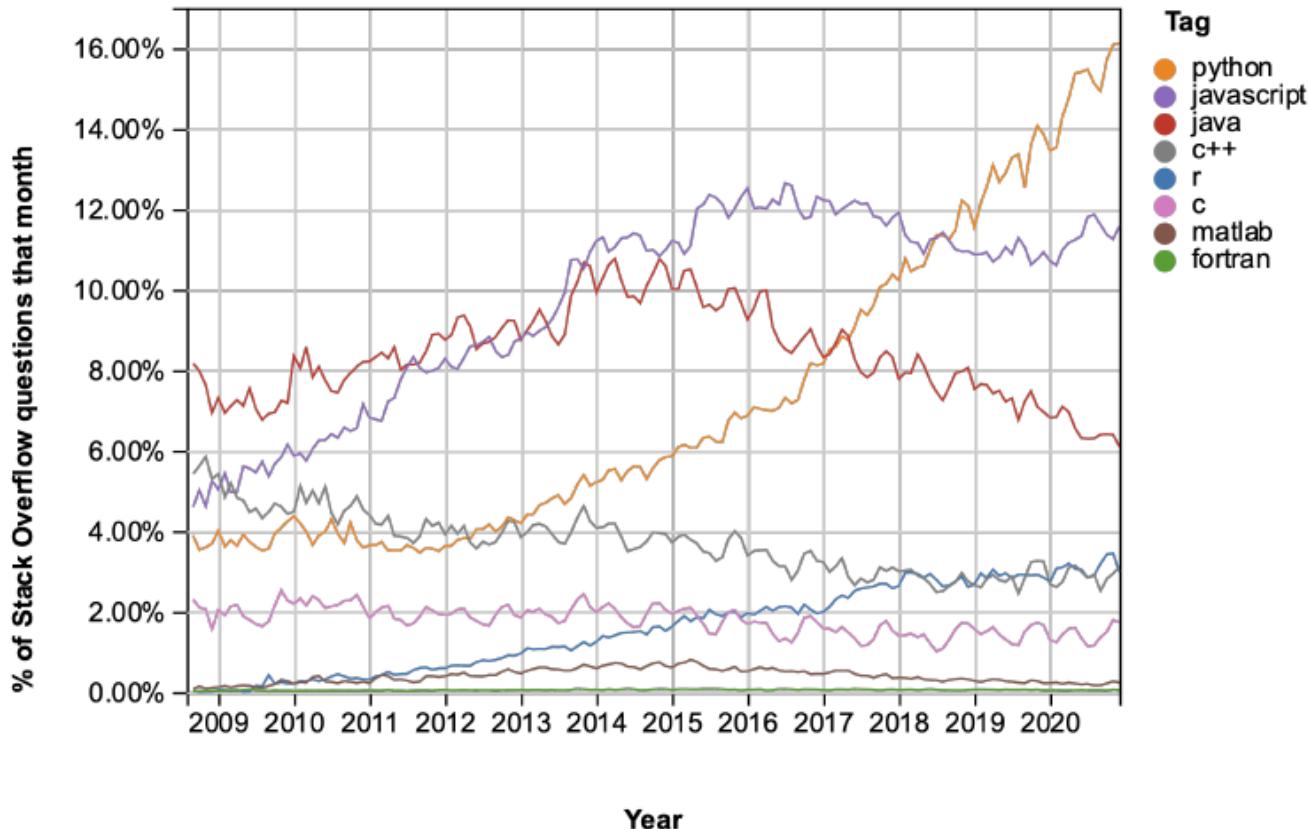


How did we get here?

1. Large, high-quality and labeled datasets
2. GPUs (graphics processing units)
3. Benchmarking competitions and algorithmic advances
4. And I argue... open source software



stack overflow trends





python

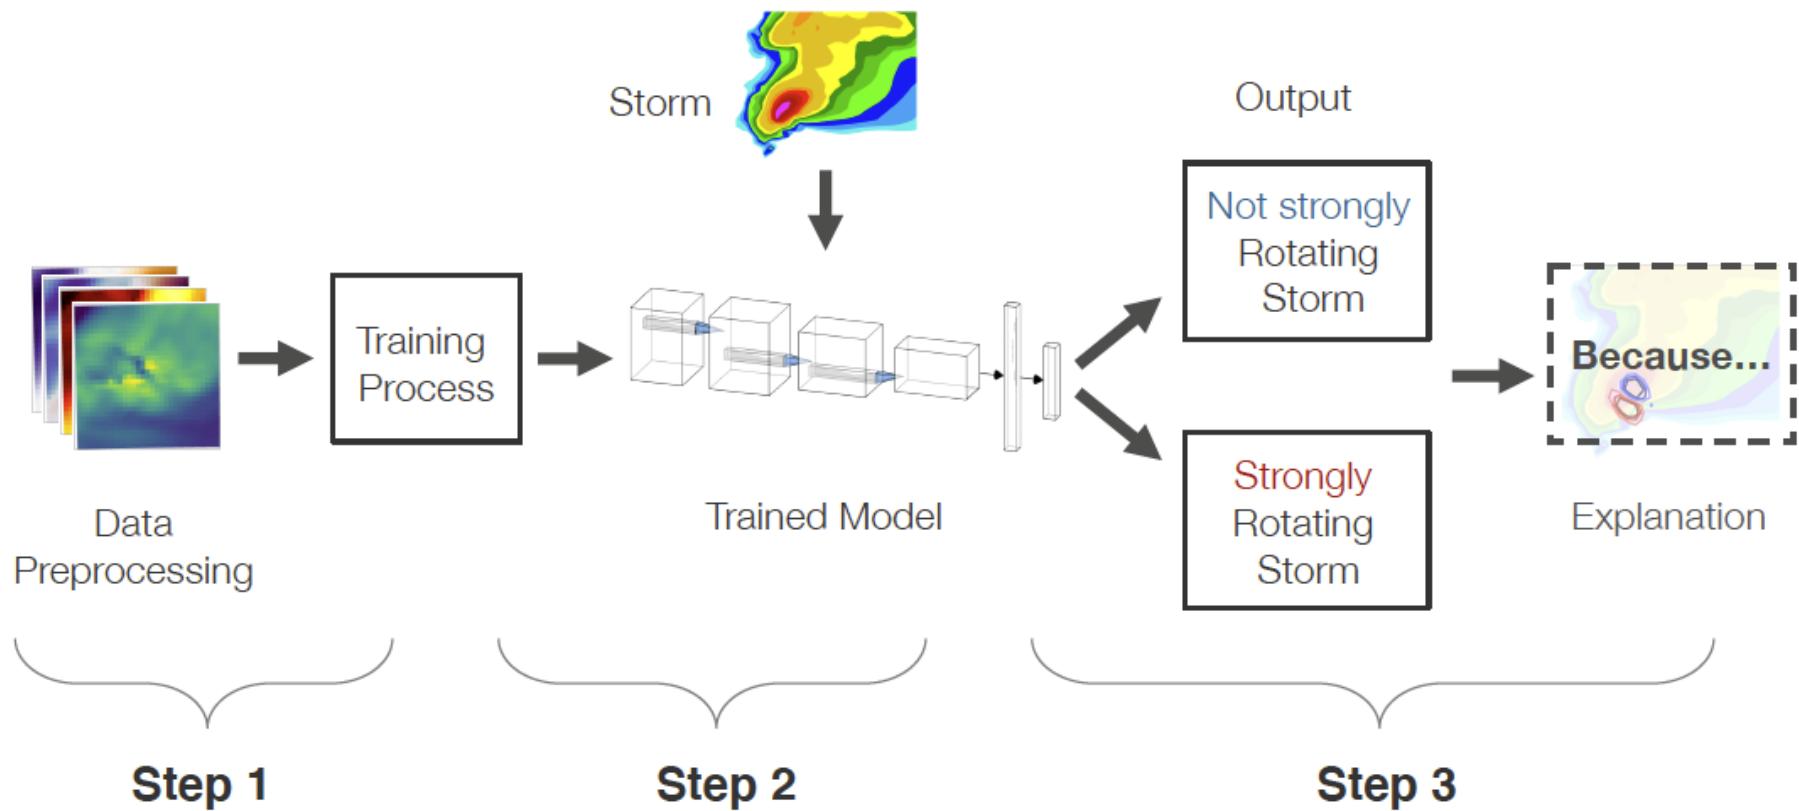
The PyTorch logo, featuring a red circle with a white outline followed by the word "PyTorch".

The Keras logo, which is a red square containing a white letter "K".

The TensorFlow logo, which is a stylized orange "F" shape.

The MXNet logo, which consists of a grid of colored squares in blue, green, and purple.

TensorFlow



Take a guess...

In your machine learning project, how much time will you spend preparing and transforming data?

- a. Less than half the time**
- b. Half the time**
- c. More than half the time**

Take a guess...

In your machine learning project, how much time will you spend preparing and transforming data?

- a. Less than half the time
- b. Half the time
- c. More than half the time

Data Availability

Data Rich

Moderate

Data Poor

Neural Networks

PINNs

Symbolic Regression

XGBoost

Transfer Learning
Meta-learning

Gaussian Processes

Use of Domain Knowledge

Adapted and Inspired from



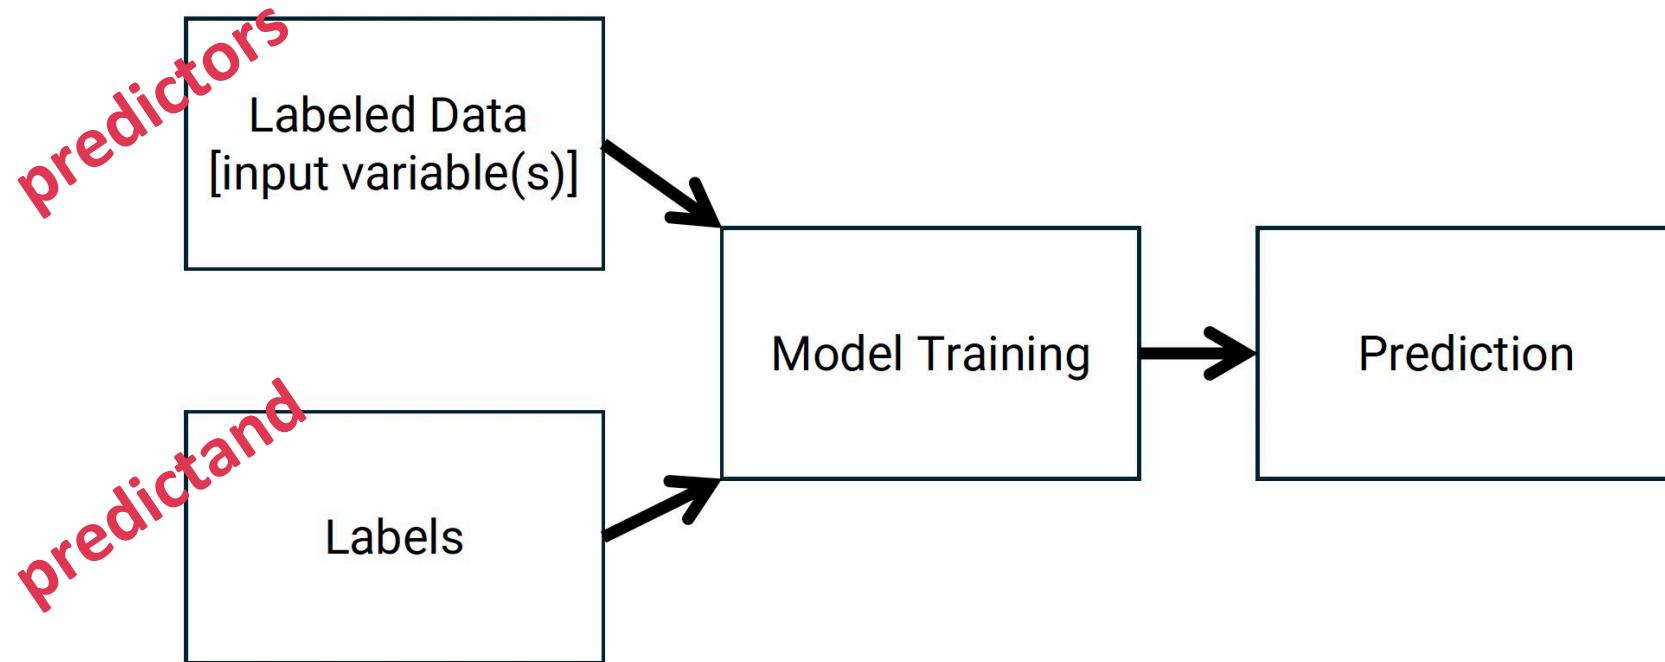
LEAP

Supervised Learning

Unsupervised Learning

Week 2	Data exploration, preprocessing, labels, ai
Week 3	Introduction to artificial neural networks
Week 4	Hyperparameter optimization of your neural networks
Week 5	Convolutional neural networks (when spatial data)
Week 6	Recurrent neural networks (when temporal data)
Week 7	Self-organizing maps
Week 8	Autoencoders

Supervised Learning: we know the answers beforehand.

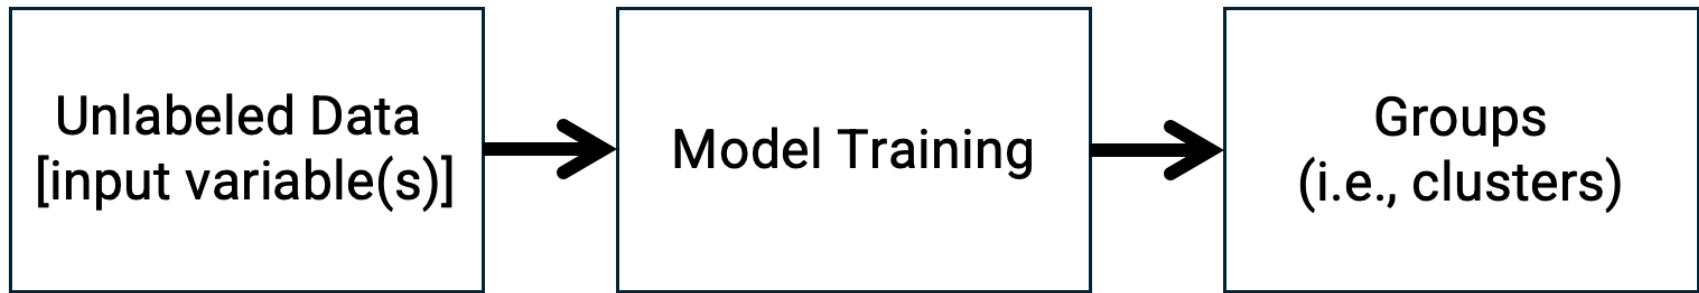


Supervised Learning: we know the answers beforehand.

$$y = mx + b$$

y are your labels; x is your input

Unsupervised Learning: we don't “know” the answers beforehand.



Unsupervised Learning: we don't "know" the answers beforehand.

$$y = mx + b$$

We have no labels (y). Usually used for grouping data or dimensionality reduction.

Semi-supervised Learning

can have unpaired data $\{x_1, x_2, \dots\}$ and labels $\{y_1, y_2, \dots\}$ and a small set of paired data $\{(x_1, y_1), (x_2, y_2) \dots\}$.

Self-supervised Learning

can sometimes be just using $\{x_1, x_2, \dots\}$ without any labels $\{y\}$, e.g., learn an useful representation of images, before training a classifier.

Nature of your “labels” (i.e., predictand or y)

Pixel (i.e., grid cell or point value)

E.g., artificial neural network (dense layers).

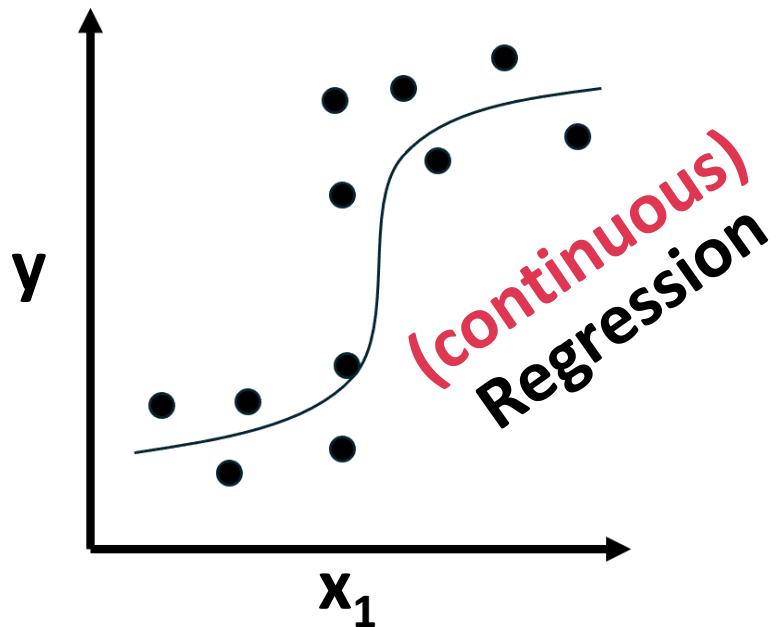
Image-based (i.e., spatial region)

E.g., convolutional neural network.

Temporal (i.e., time series)

E.g., recurrent neural network.

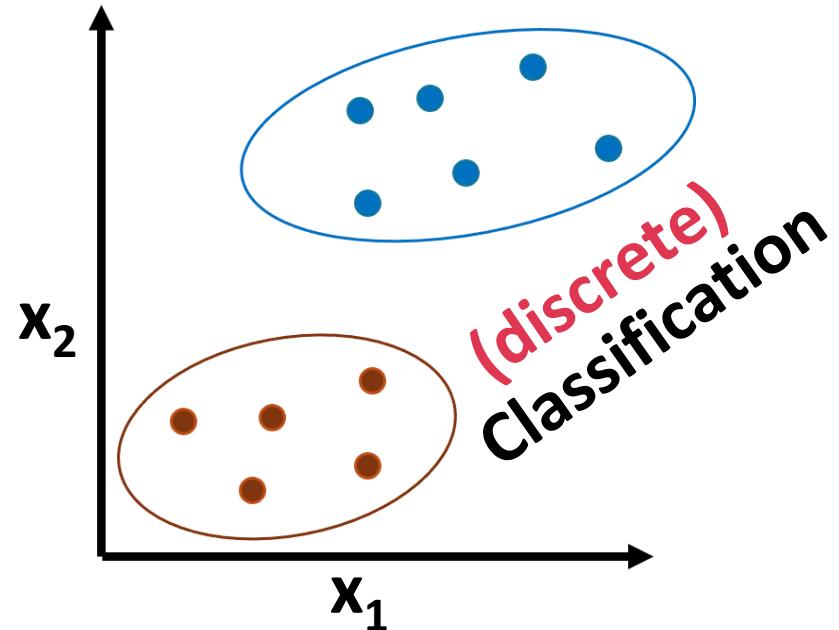
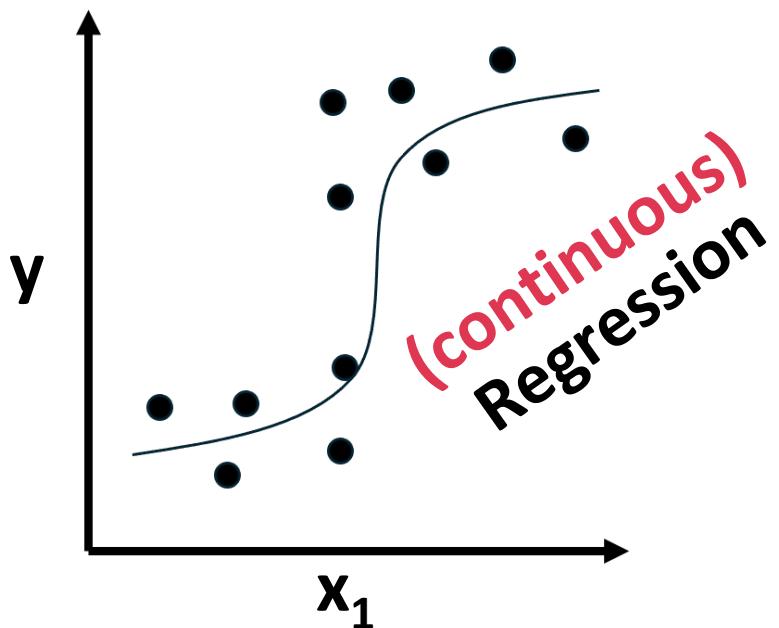
Predict the *value* of something.



Some examples of regression problems include:

- tropical cyclone maximum wind speed at a certain lead time.
- maximum land surface temperature in an urban area.
- percentage of cloud coverage over a predefined domain.
- bias correction of subseasonal forecasts of precipitation.

Predict the *value* of something.



Predict the *class* of something.

Examples of classification problems:

- whether a thunderstorm will produce severe hazards or be non-severe (i.e., binary).
- whether a polynya is occurring in the Ross Sea (i.e., binary).
- type of cloud(s) that is(are) in a citizen science photograph (i.e., multiclass).
- subseasonal prediction of a North American weather regime class (i.e., multiclass).

Take a guess...

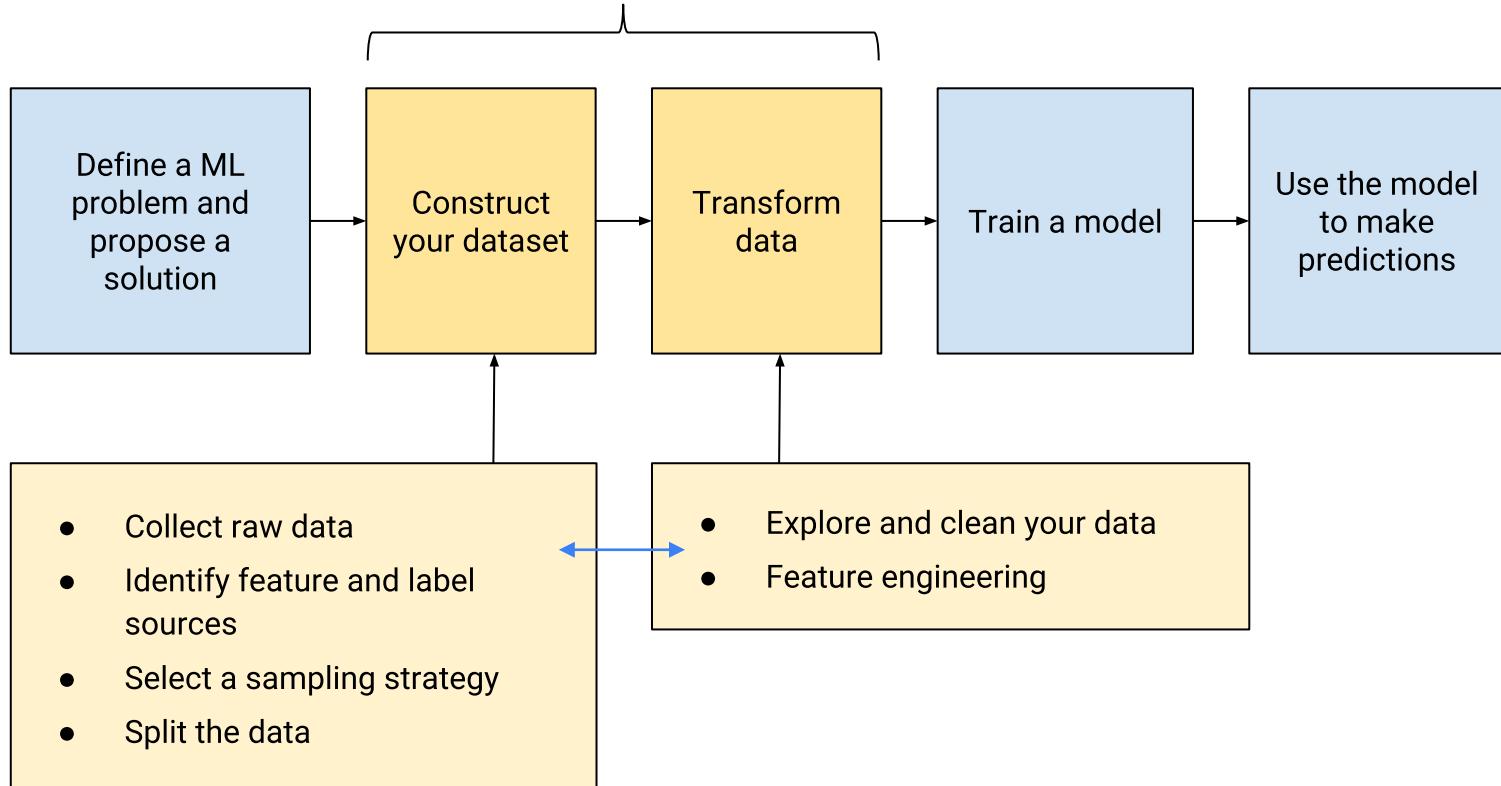
If you had to prioritize one of the areas below for your ML project, which one would have the most impact?

- a. Using the latest optimization algorithm
- b. Quality and size of your training data
- c. Deeper network
- d. More clever loss function

Take a guess...

If you had to prioritize one of the areas below for your ML project, which one would have the most impact?

- a. Using the latest optimization algorithm
- b. Quality and size of your training data**
- c. Deeper network
- d. More clever loss function



Questions to ask yourself...

- How many features (i.e., variables) should I use?
- Is my data reliable, representative, or skewed?
- Will I use direct labels or derived (proxy) labels?
- Do I have too much data or imbalanced data?
- How should I split my data into train and test sets?
- How and when should I transform my data?
- How do I encode my classes?

So what can you do about this class issue?

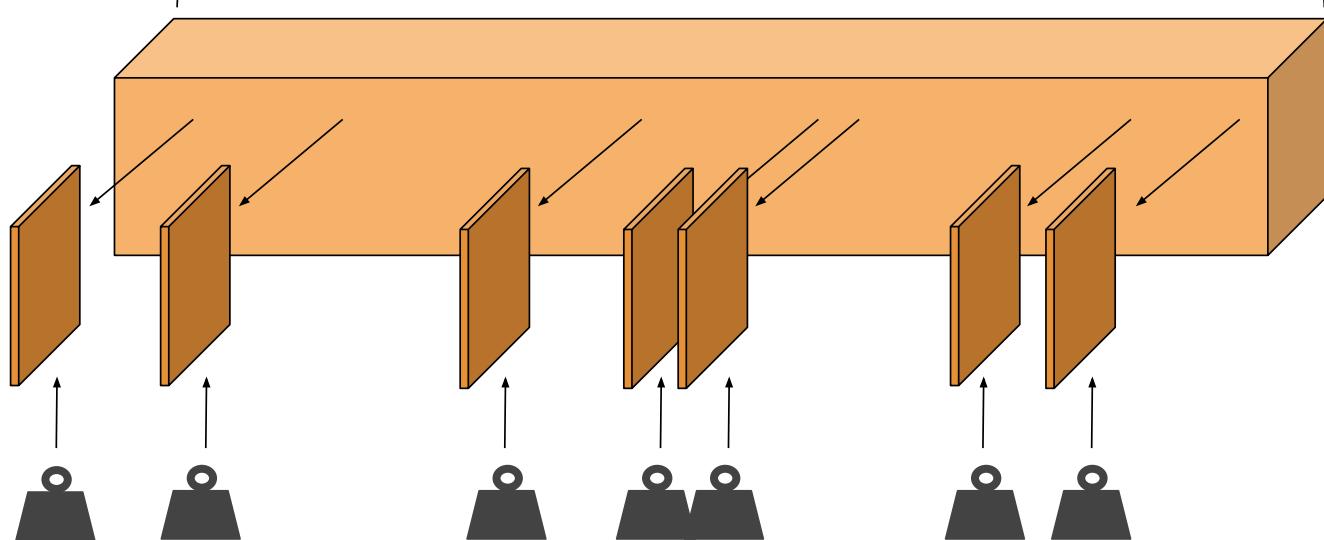
Downsampling and Upweighting

- **Downsampling** (in this context) means training on a disproportionately low subset of the majority class examples.
- **Adding minority class weights** means adding example weight(s) to the minority class.

Dataset of dominant class examples

1 Downsample

Extract random examples from the dominant class.



2 Upweight

Add weight to the downsampled examples.

$$\{\text{example weight}\} = \{\text{original example weight}\} \times \{\text{downsampling factor}\}$$

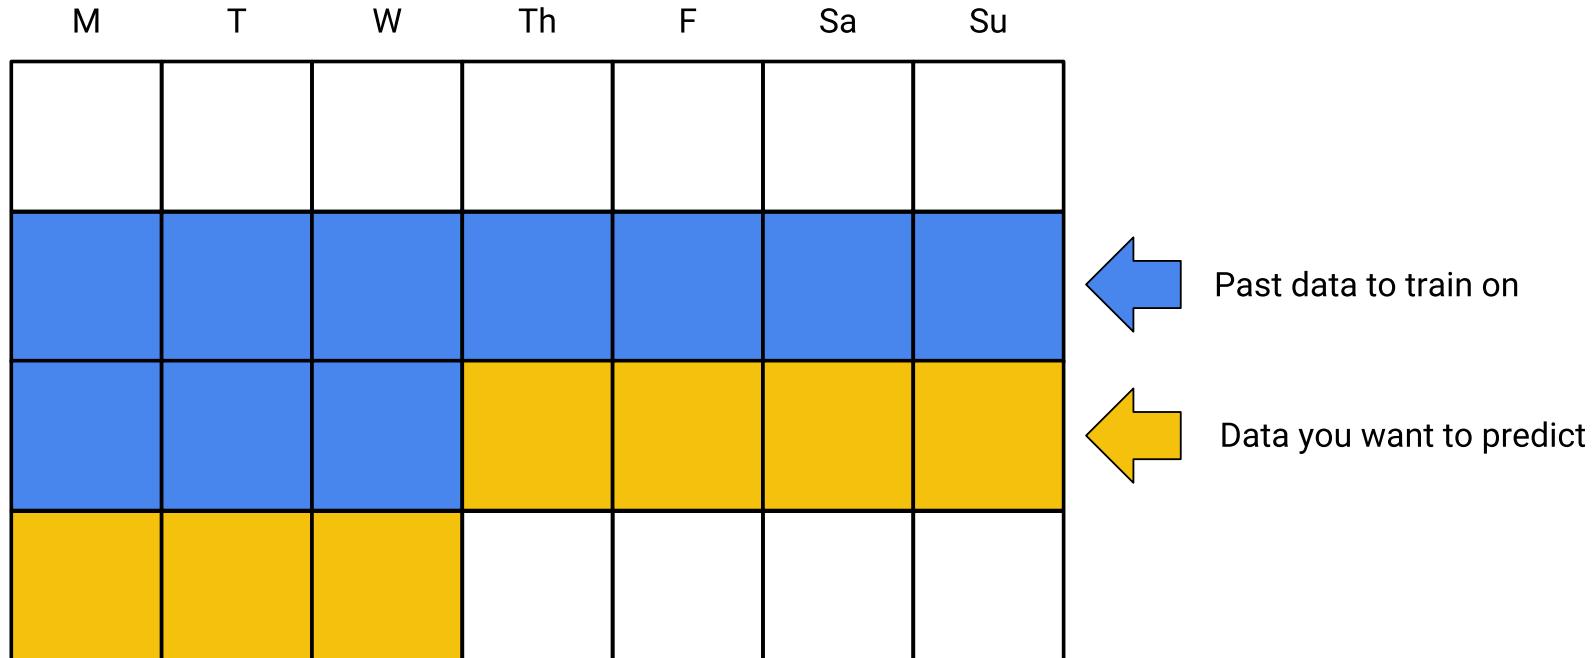
Faster convergence: During training, we see the minority class more often, which will help converge faster.

Disk space: Consolidating the majority class into fewer examples with larger weights, we spend less disk space storing them.

Calibration: Upweighting ensures our model is still calibrated; the outputs can still be interpreted as probabilities.

Seasonality, trends, or cyclical effects can be an issue!

Consider removing these effect during preprocessing or using the date as a feature so the model can learn these effects.

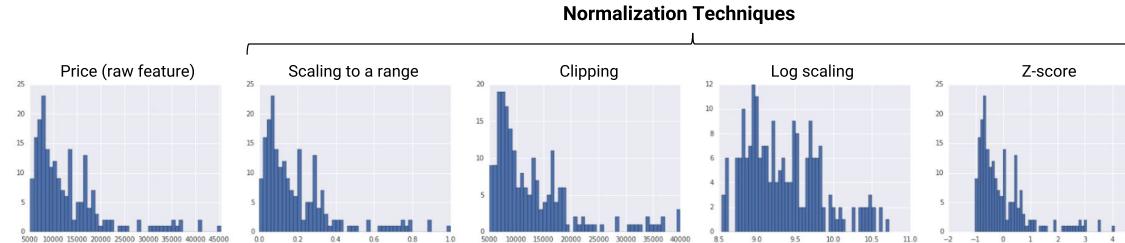


Where/when to do transformations?

Transformations can be done before or during training.

Good when data statistics are needed. Dangerous during inference. Slow when iterating.

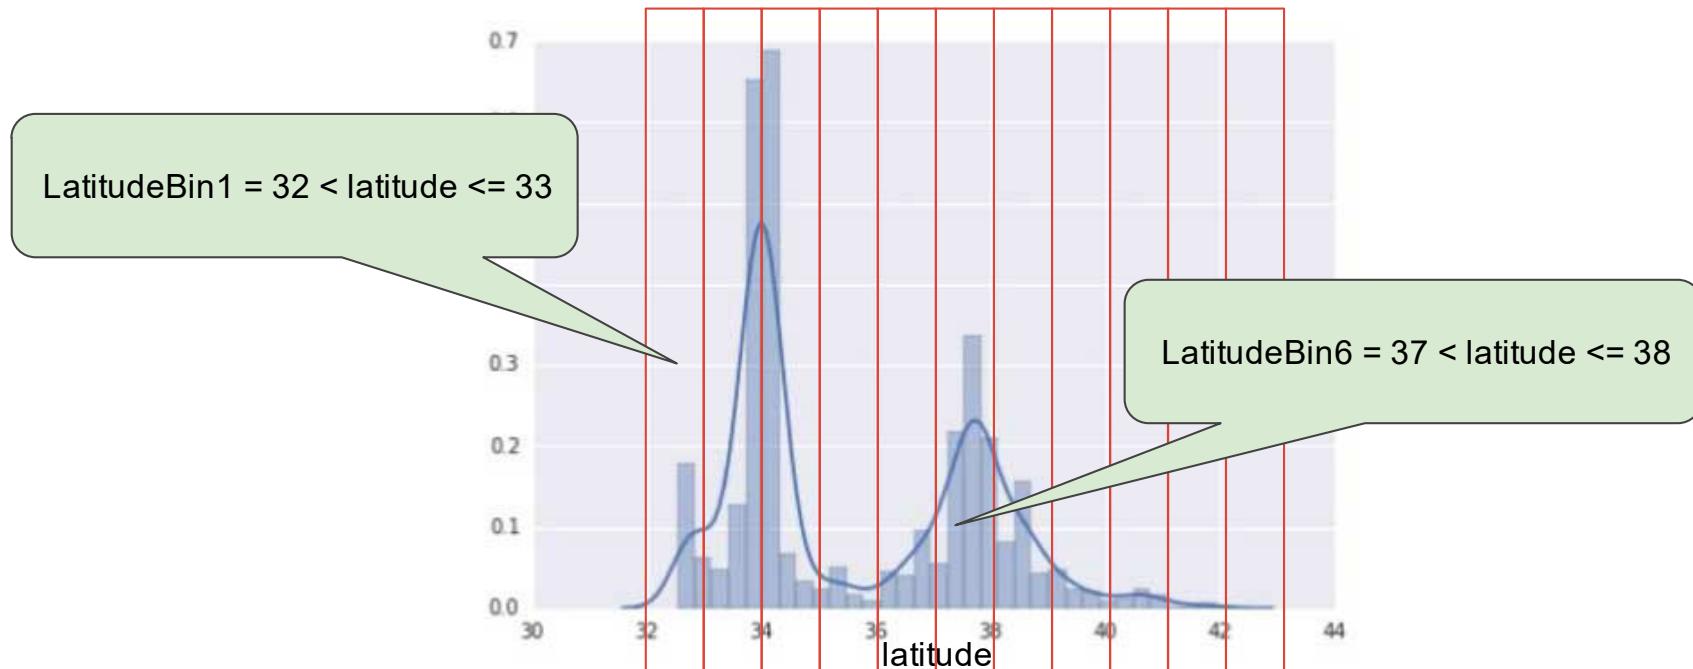
Easy and fast to iterate. Transforms per batch can slow training if complicated.



We can frame our problem to be **binary** classification (two classes; e.g., True or False, or Yes or No problems).

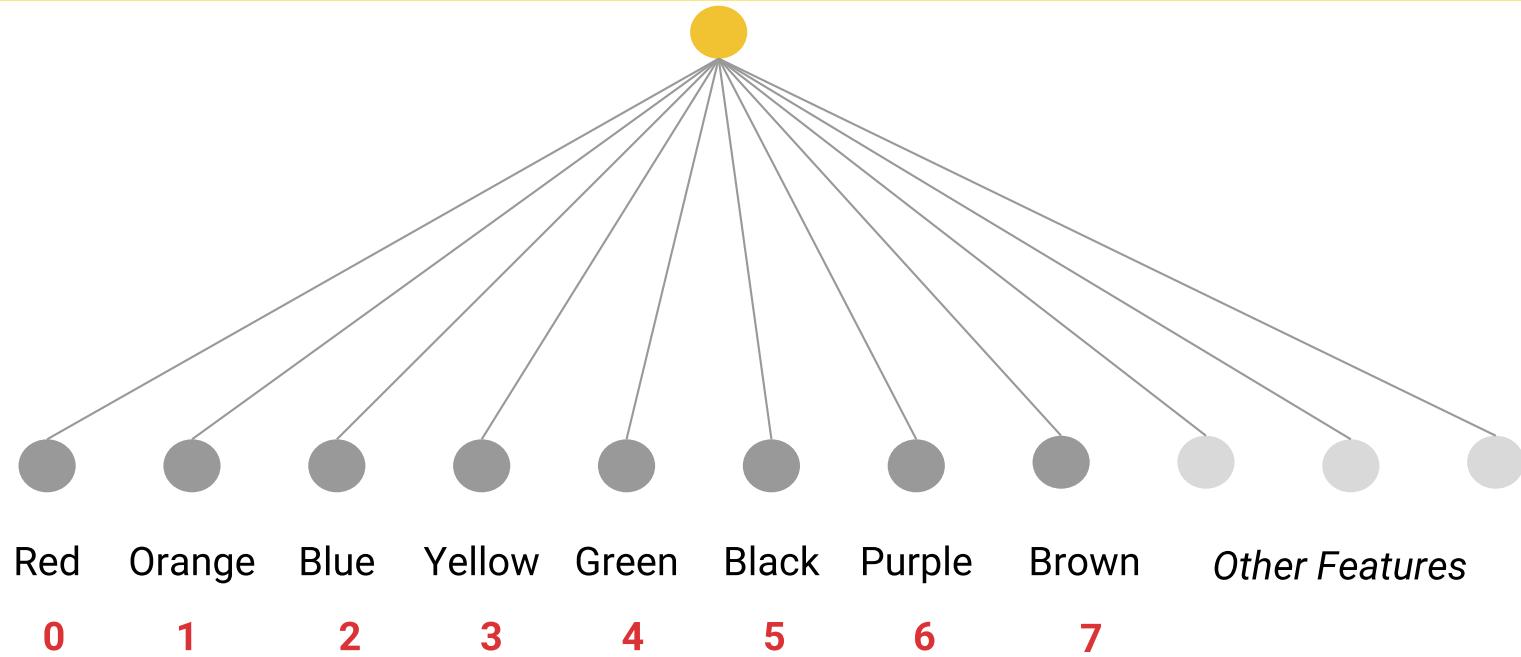
Classification problems can also be **multiclass** in nature.

Continuous data could be grouped or binned into certain classes (i.e., buckets); *continuous data can be converted to classes.*

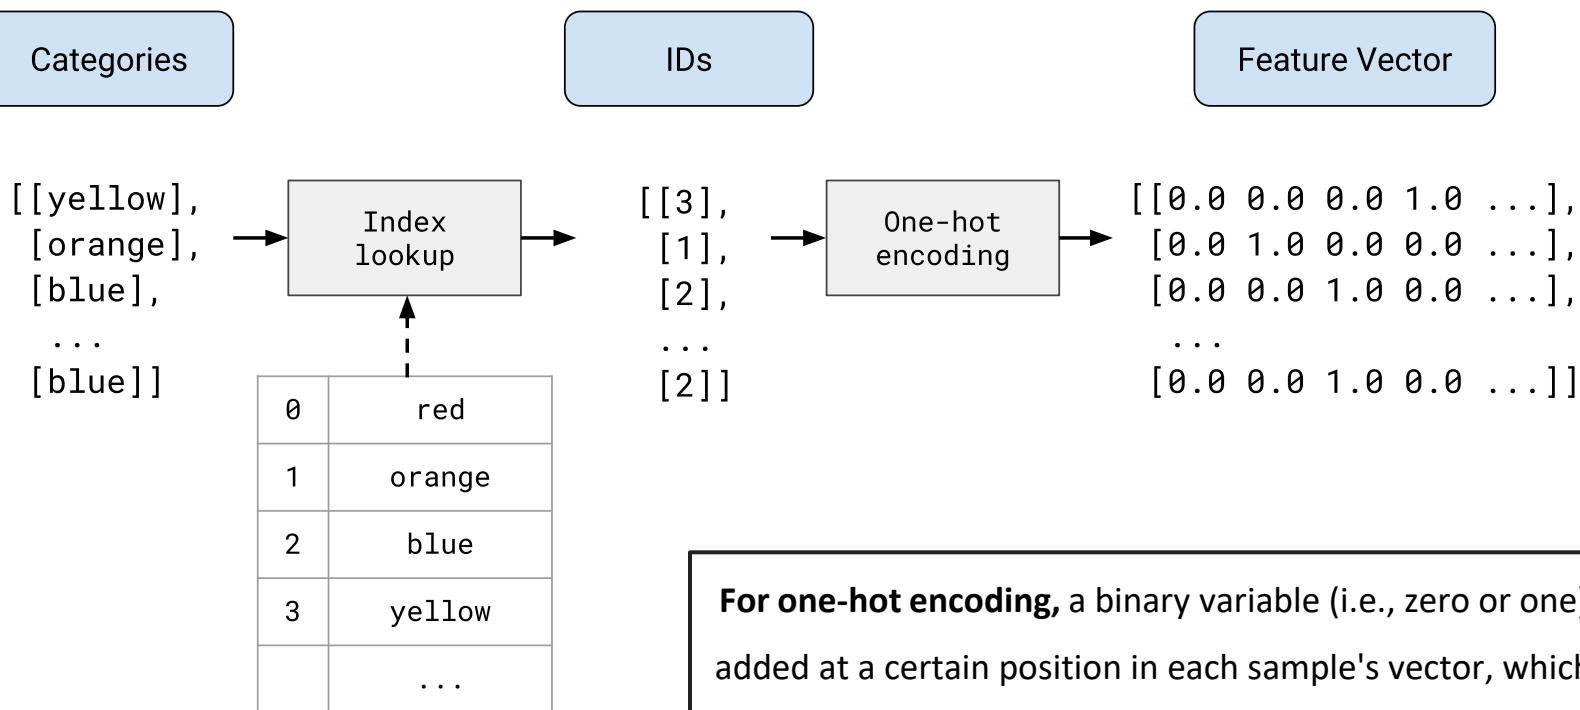


An example of a common occurrence of binning:

- Bin continuous maximum tropical cyclone wind speeds into hurricane intensities based on the Saffir-Simpson Hurricane Wind Scale, e.g., category 1, category 2, etc.



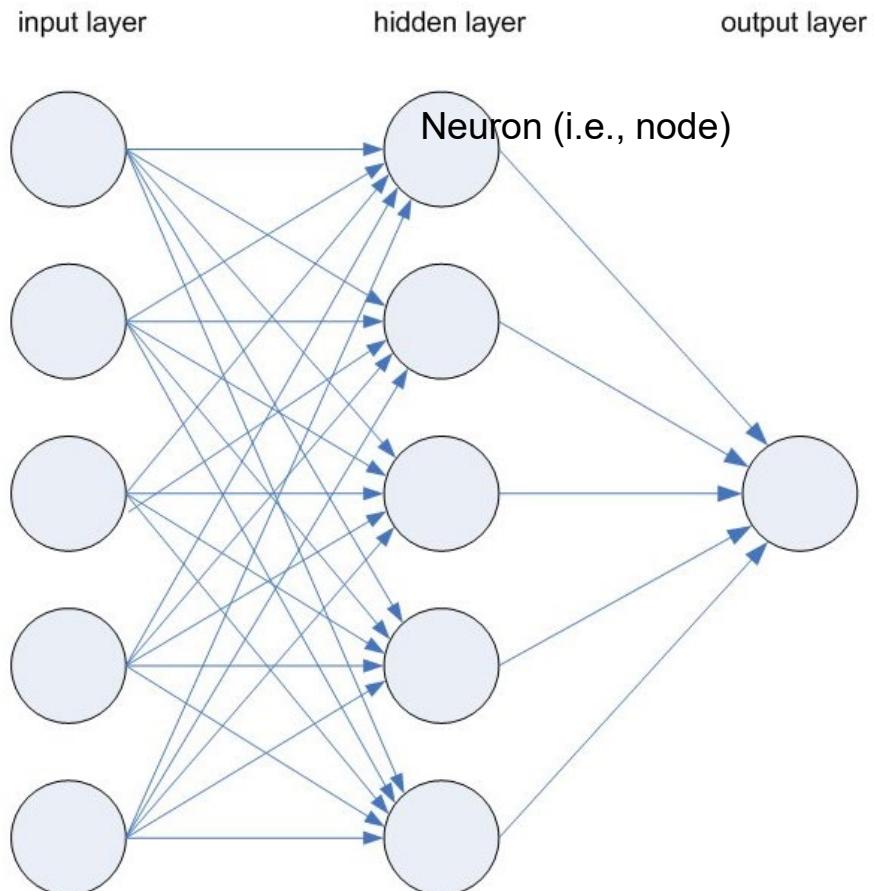
Integer encoding involves assigning an integer value to each unique category. This can work well (not mandatory) for data that have an underlying natural ordered relationship between each other, since machine learning models may be able to understand and harness this relationship.

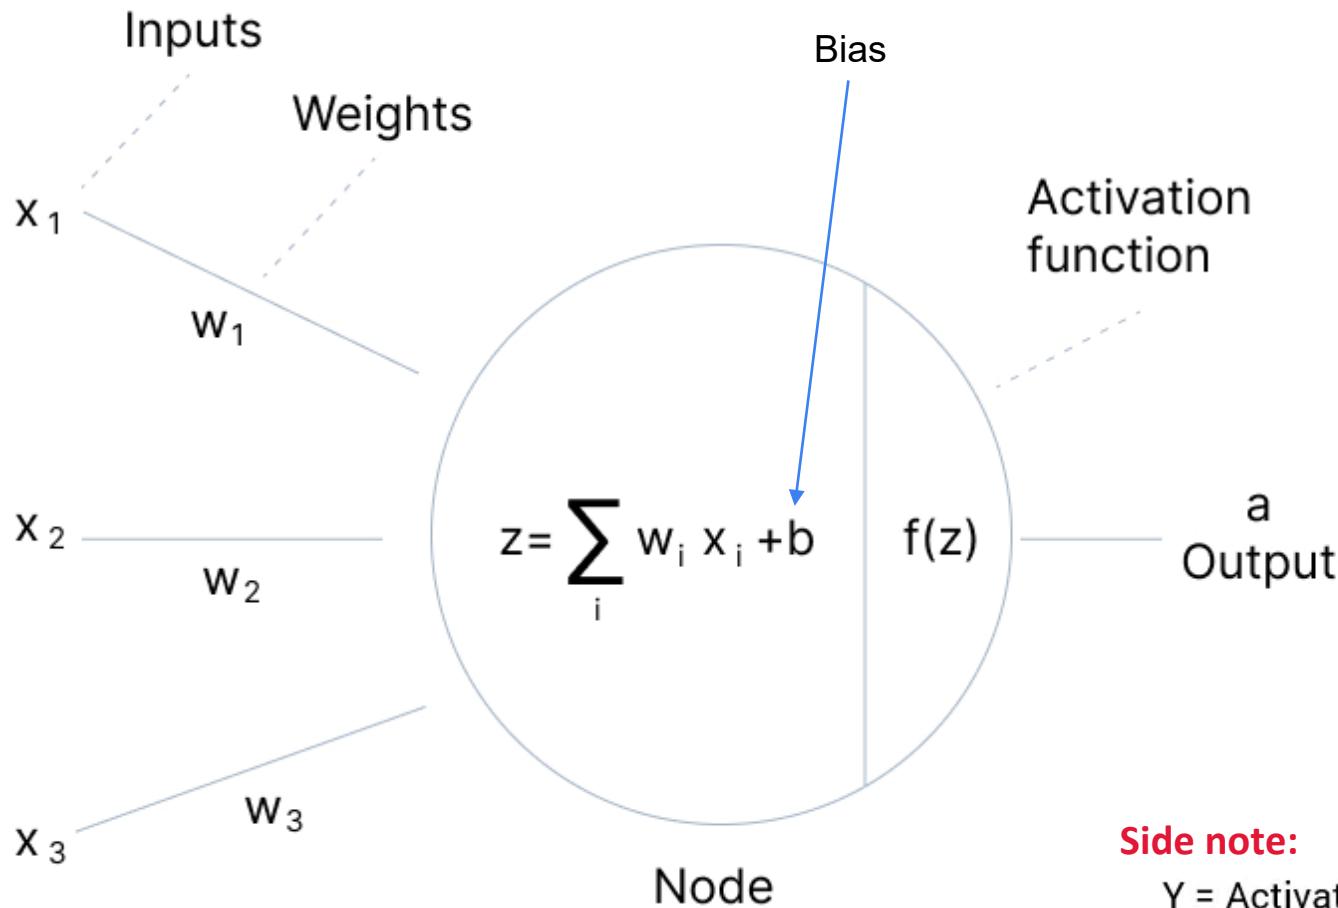


For one-hot encoding, a binary variable (i.e., zero or one) is added at a certain position in each sample's vector, which is meant to represent the respective sample's class, e.g., [0, 0, 1] vs [1, 0, 0].

Neural networks are machine learning models consisting of layers and neurons (i.e., nodes).

The basic units are neurons (i.e., nodes), which are typically organized into layers.



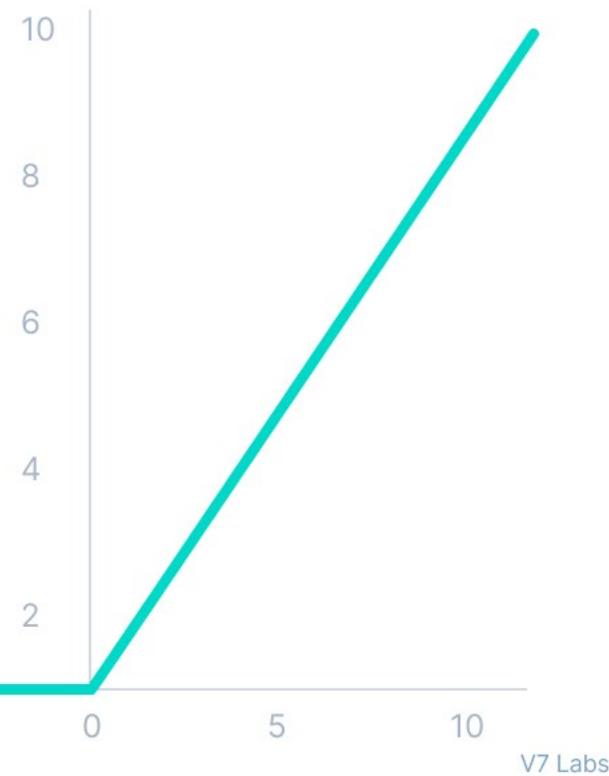


Weights and bias terms are learned. Activation function is not; it is a "hyperparameter."

Side note:

$$Y = \text{Activation}(\sum(\text{weight} * \text{input}) + \text{bias})$$

ReLU



$$f(x) = \max(0, x)$$

Since ReLU outputs zero for all negative inputs, it naturally leads to sparse activations; at any time, only a subset of neurons are activated, leading to more efficient computation.

Hyperparameters are parameters whose values are set before starting the model training process.

Method	Description
Learning rate	Determines the rate update/step size for weights during backpropagation. Usually between 0.01 to 0.0001.
Activation functions	Function that calculates the output of a node based on the sum of the inputs times weights and bias. Can add nonlinearity.
Weight initialization	Method to initialize the weights. Can be zero, one, random...
Preprocessing choices	How you preprocess your data is an important hyperparameter!
Loss function	What error would you like to minimize? E.g., MSE, RMSE, etc.

Hyperparameters are parameters whose values are set before starting the model training process.

Method	Description
Epoch	How many times the entire training dataset can be “seen” by the neural network. Too many epochs can lead to overfitting.
Batch size	How many samples from the training set can be “seen” at one time by the neural network to compute error with and perform backprop. Usually numbers like 16, 32, 64, 128, or 256.
Train and test split	How you split your dataset can impact neural network performance. 80/20 or 60/40 splits are common.
Layer types	Types of layers (e.g., dense, fully connected).

That's it for me!

Over to Tom Beucler now!

3) How are ML models trained & evaluated?

4) What is physics-guided ML?

That's it for me!

Over to Tom Beucler now!

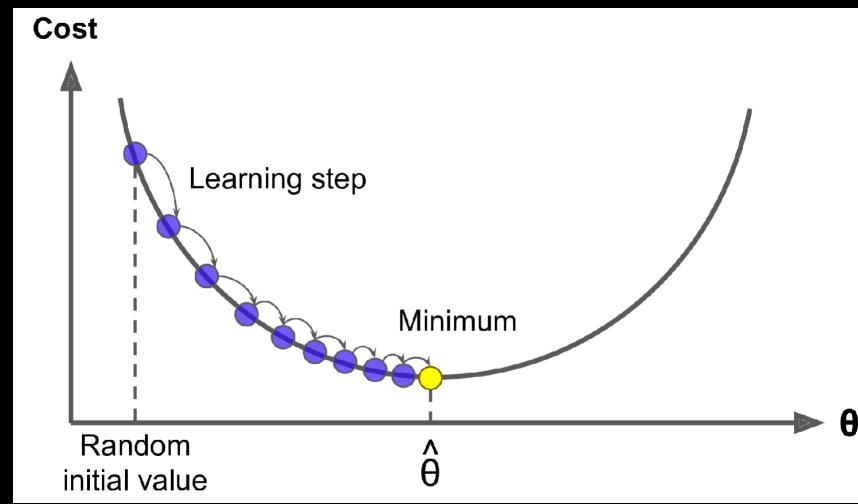
3) How are ML models trained & evaluated?

4) What is physics-guided ML?

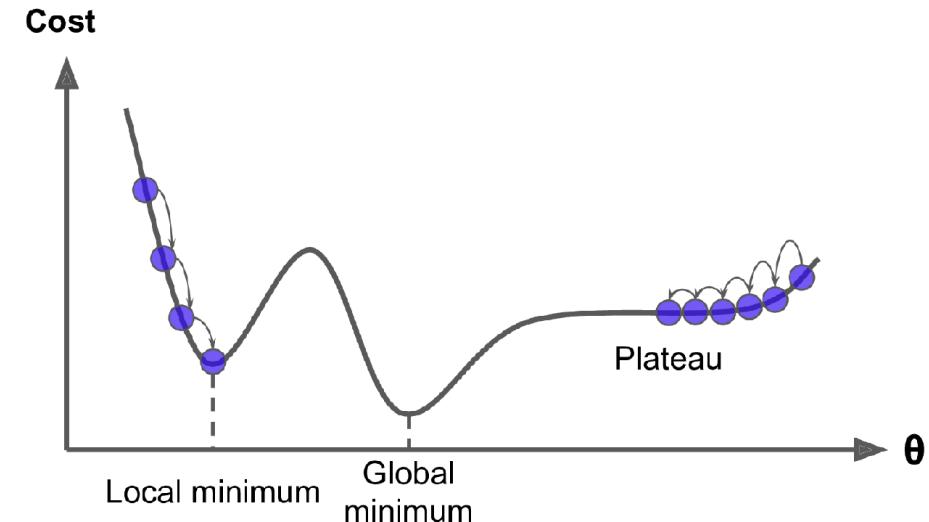
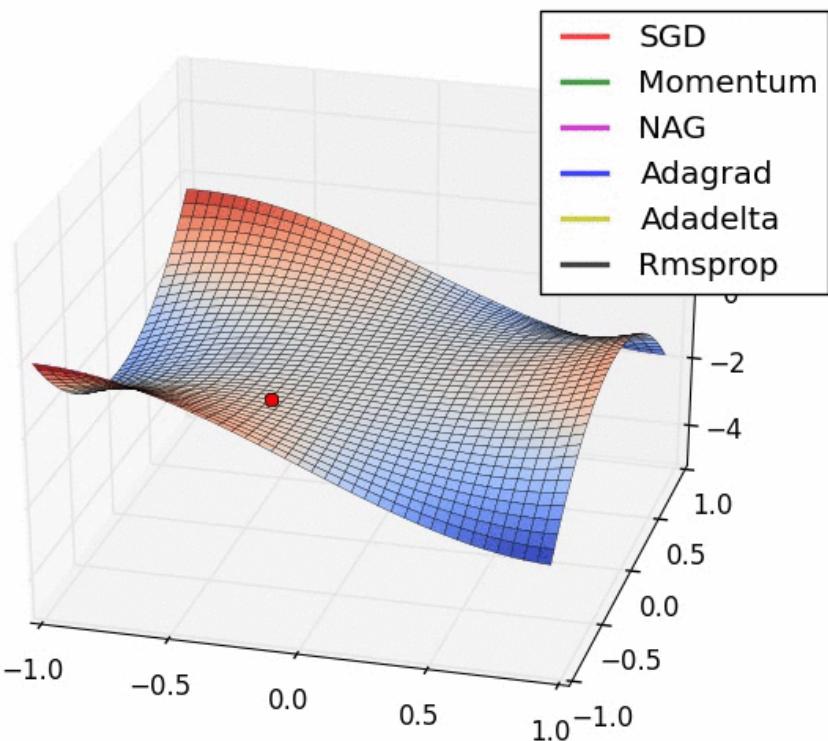
1) What is a cost/loss/error function?

Function: (Predictions, [Truth], [Inputs]) \mapsto (Cost/Loss/Error to minimize)

For complex data, usually minimized using gradient descent:



1) Training ML Models: Gradient Descent



Batch vs Stochastic vs Mini-Batch
Gradient Descent

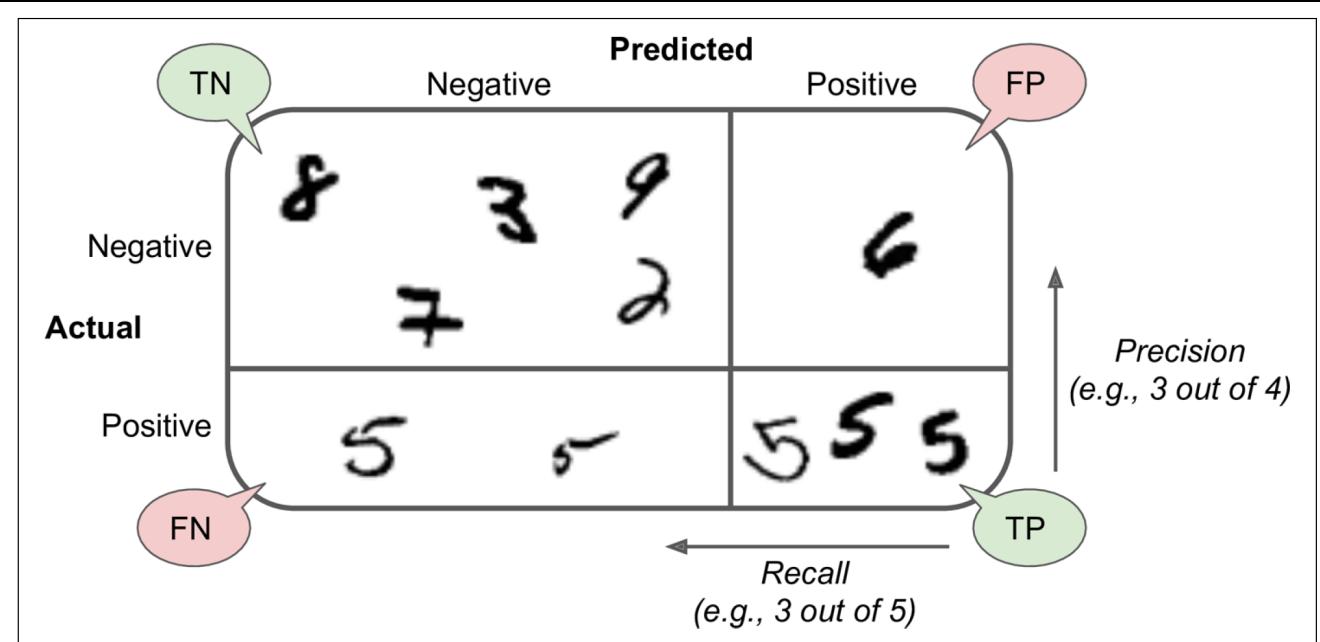
Source: Computational Optimization (Cornell University)

TROPICANA ML Tutorial

2) Many options to train (and evaluate) classifiers!

Classification metrics	
See the Classification metrics section of the user guide for further details.	
<code>metrics.accuracy_score(y_true, y_pred, *[...])</code>	Accuracy classification score.
<code>metrics.auc(x, y)</code>	Compute Area Under the Curve (AUC) using the trapezoidal rule.
<code>metrics.average_precision_score(y_true, ...)</code>	Compute average precision (AP) from prediction scores.
<code>metrics.balanced_accuracy_score(y_true, ...)</code>	Compute the balanced accuracy.
<code>metrics.brier_score_loss(y_true, y_prob, *)</code>	Compute the Brier score loss.
<code>metrics.classification_report(y_true, y_pred, *)</code>	Build a text report showing the main classification metrics.
<code>metrics.cohen_kappa_score(y1, y2, *[...])</code>	Cohen's kappa: a statistic that measures inter-annotator agreement.
<code>metrics.confusion_matrix(y_true, y_pred, *)</code>	Compute confusion matrix to evaluate the accuracy of a classification.
<code>metrics.dcg_score(y_true, y_score, *[t, k, ...])</code>	Compute Discounted Cumulative Gain.
<code>metrics.det_curve(y_true, y_score[, ...])</code>	Compute error rates for different probability thresholds.
<code>metrics.f1_score(y_true, y_pred, *[..., beta])</code>	Compute the F1 score, also known as balanced F-score or F-measure.
<code>metrics.fbeta_score(y_true, y_pred, *, beta)</code>	Compute the F-beta score.
<code>metrics.hamming_loss(y_true, y_pred, *[..., ...])</code>	Compute the average Hamming loss.
<code>metrics.hinge_loss(y_true, pred_decision, *)</code>	Average hinge loss (non-regularized).
<code>metrics.jaccard_score(y_true, y_pred, *[..., ...])</code>	Jaccard similarity coefficient score.
<code>metrics.log_loss(y_true, y_pred, *[..., eps, ...])</code>	Log loss, aka logistic loss or cross-entropy loss.
<code>metrics.matthews_corrcoef(y_true, y_pred, *)</code>	Compute the Matthews correlation coefficient (MCC).
<code>metrics.multilabel_confusion_matrix(y_true, ...)</code>	Compute a confusion matrix for each class or sample.
<code>metrics.ndcg_score(y_true, y_score, *[..., k, ...])</code>	Compute Normalized Discounted Cumulative Gain.
<code>metrics.precision_recall_curve(y_true, ...)</code>	Compute precision-recall pairs for different probability thresholds.
<code>metrics.precision_recall_fscore_support(...)</code>	Compute precision, recall, F-measure and support for each class.
<code>metrics.precision_score(y_true, y_pred, *[..., ...])</code>	Compute the precision.
<code>metrics.recall_score(y_true, y_pred, *[..., ...])</code>	Compute the recall.
<code>metrics.roc_auc_score(y_true, y_score, *[..., ...])</code>	Compute Area Under the Receiver Operating Characteristic Curve (ROC AUC) from prediction scores.
<code>metrics.roc_curve(y_true, y_score, *[..., ...])</code>	Compute Receiver operating characteristic (ROC).
<code>metrics.top_k_accuracy_score(y_true, y_score, *)</code>	Top-k Accuracy classification score.
<code>metrics.zero_one_loss(y_true, y_pred, *[..., ...])</code>	Zero-one classification loss.

2) Confusion Matrix of the task: Is the digit 5?

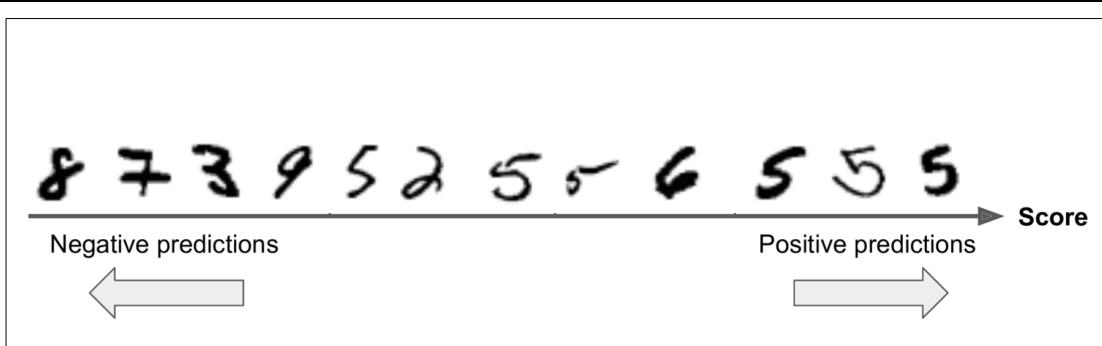


Accuracy:

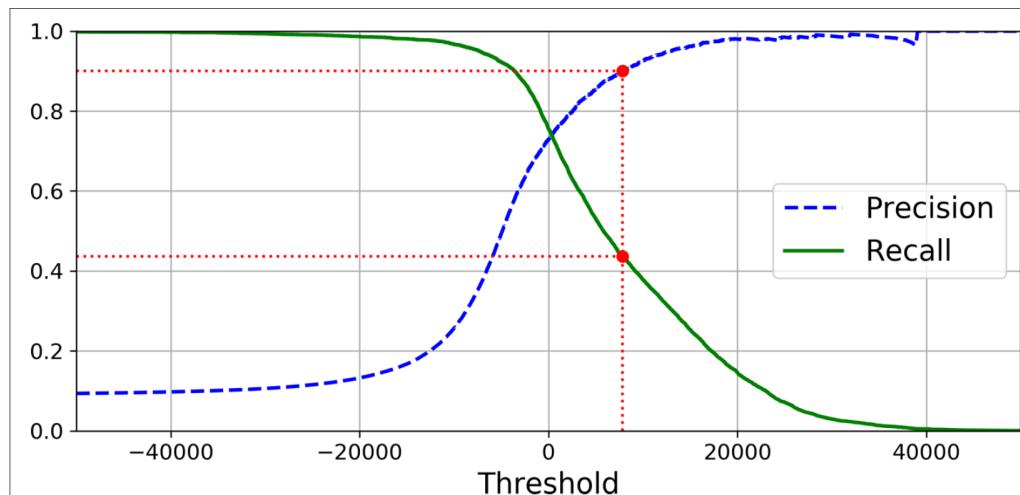
$$\frac{TP+TN}{TP+TN+FP+FN}$$

Figure 3-2. An illustrated confusion matrix shows examples of true negatives (top left), false positives (top right), false negatives (lower left), and true positives (lower right)

TROPICANA ML Tutorial



$$\text{recall} = \frac{TP}{TP + FN}$$



$$\text{precision} = \frac{TP}{TP + FP}$$

Figure 3-4. Precision and recall versus the decision threshold

3) Loss/Cost/Error Functions for Regression Models

Mean Squared Error:

$$\text{MSE} = \frac{1}{n} \sum_{i=1}^n (Y_i - \hat{Y}_i)^2$$

For environmental applications:

- Root-Mean-Square Error (RMSE)
- Mean Absolute Error

Are popular because they preserve units!



“Truth” Pred.

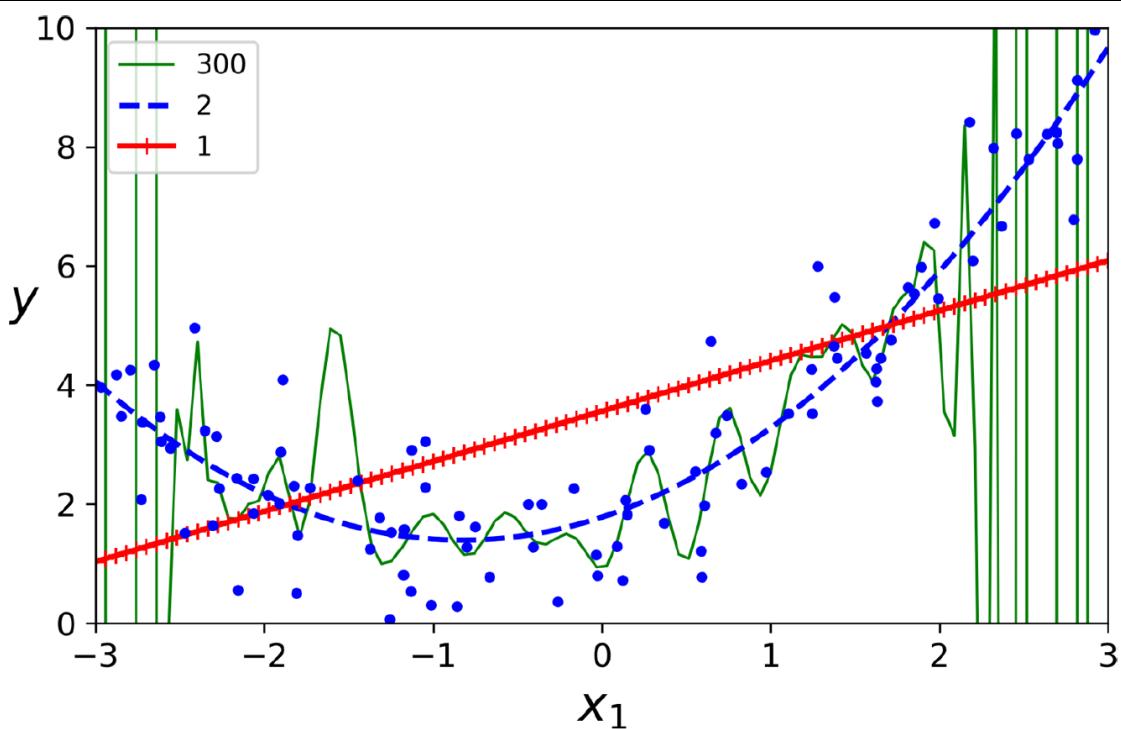
3) Many options to train (and evaluate) regression models!

Regression metrics

See the [Regression metrics](#) section of the user guide for further details.

<code>metrics.explained_variance_score(y_true, ...)</code>	Explained variance regression score function.
<code>metrics.max_error(y_true, y_pred)</code>	The max_error metric calculates the maximum residual error.
<code>metrics.mean_absolute_error(y_true, y_pred, *)</code>	Mean absolute error regression loss.
<code>metrics.mean_squared_error(y_true, y_pred, *)</code>	Mean squared error regression loss.
<code>metrics.mean_squared_log_error(y_true, y_pred, *)</code>	Mean squared logarithmic error regression loss.
<code>metrics.median_absolute_error(y_true, y_pred, *)</code>	Median absolute error regression loss.
<code>metrics.mean_absolute_percentage_error(...)</code>	Mean absolute percentage error (MAPE) regression loss.
<code>metrics.r2_score(y_true, y_pred, *[...])</code>	R^2 (coefficient of determination) regression score function.
<code>metrics.mean_poisson_deviance(y_true, y_pred, *)</code>	Mean Poisson deviance regression loss.
<code>metrics.mean_gamma_deviance(y_true, y_pred, *)</code>	Mean Gamma deviance regression loss.
<code>metrics.mean_tweedie_deviance(y_true, y_pred, *)</code>	Mean Tweedie deviance regression loss.
<code>metrics.d2_tweedie_score(y_true, y_pred, *)</code>	D^2 regression score function, percentage of Tweedie deviance explained.
<code>metrics.mean_pinball_loss(y_true, y_pred, *)</code>	Pinball loss for quantile regression.

How can ML best practices help us consistently assess the performance of our three models?

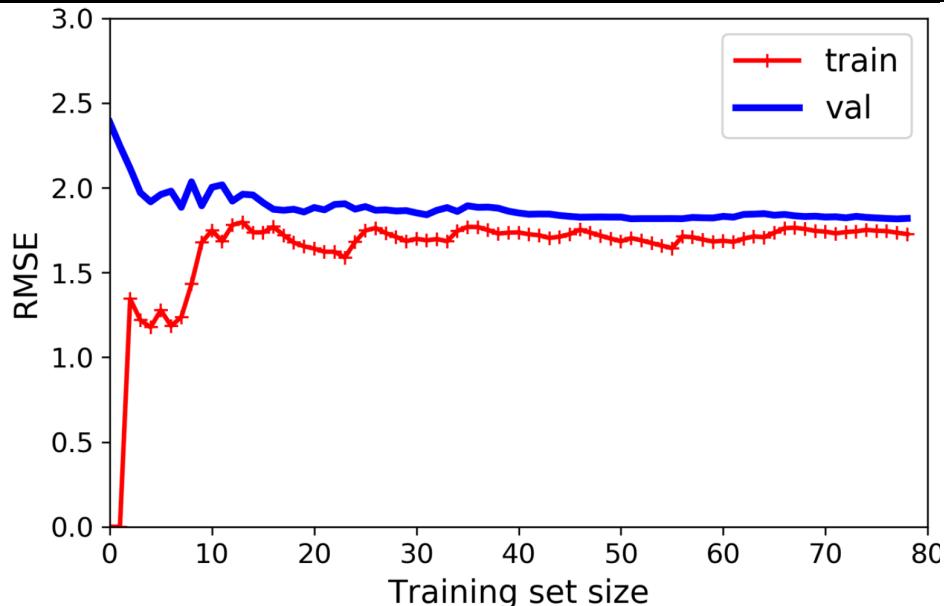


Which model would you choose in practice? Why?

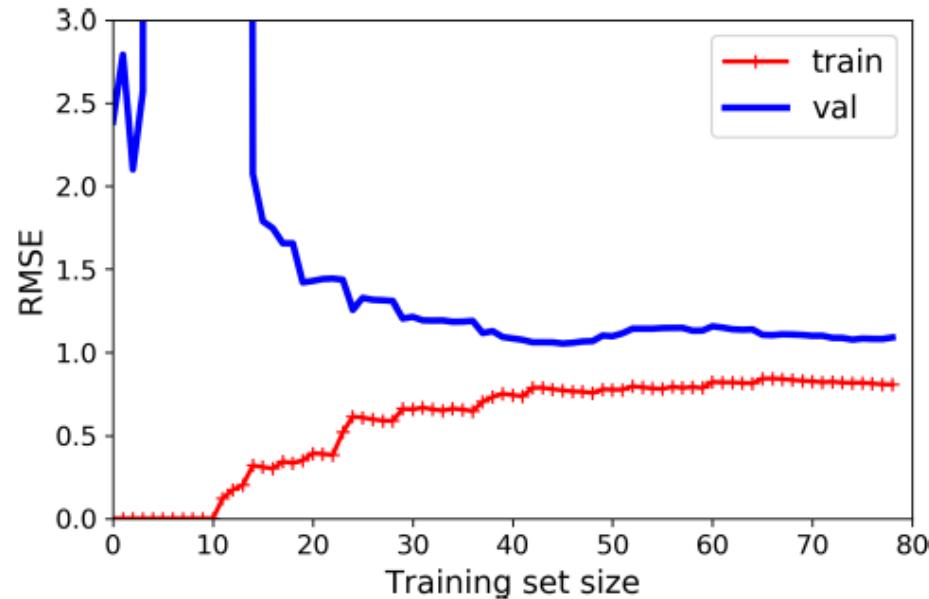
The quadratic model because the polynomial regression model is overfitting the training data and the linear model is underfitting the training data.

Cross-Validation (& Learning Curves) Help Assess how ML Models Generalize to Out-of-Sample Data

Linear Regression: Underfitting



High-degree Polynomial: Overfitting



How to benchmark a regression model?

Step 1: Pick performance metrics (MSE, RMSE, MAE, etc.)

Coefficient of **determination** (between $-\infty$ and 1):

$$R^2 = 1 - \frac{SS_{\text{res}}}{SS_{\text{tot}}}$$

← Sum of squared errors (proportional to MSE)
← Total Sum of squares (proportional to true variance)



≠ Coefficient of **correlation**:
(between -1 and 1)

$$\frac{\text{cov}(Y_{\text{true}}, Y_{\text{pred}})}{\text{std}(Y_{\text{true}}) \times \text{std}(Y_{\text{pred}})}$$

How to benchmark a regression model?

Step 2: Report mean performance over training, validation, and test sets

	Training Set	Validation Set	Test Set
RMSE (kg)	10	11	50
Coefficient of Determination	0.6	0.45	-2

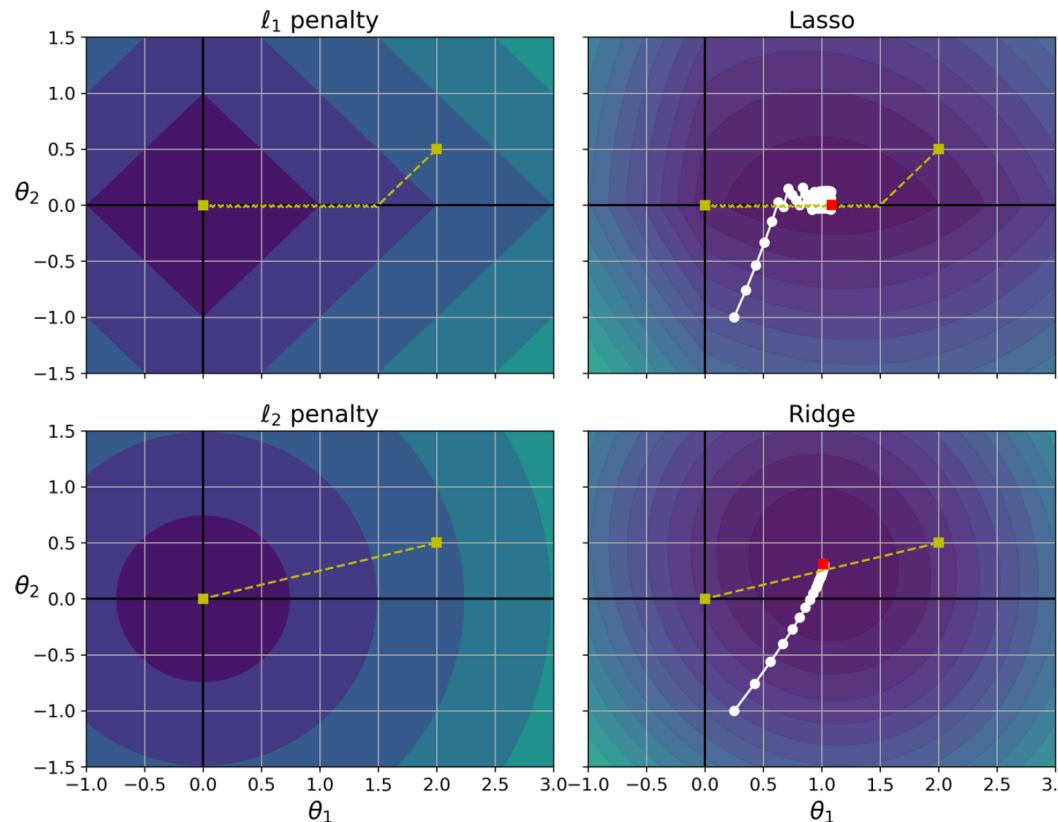


Figure 4-19. Lasso versus Ridge regularization

Lasso: Penalizes the absolute value of the weights (L1 norm)

≠

Ridge: Penalizes the squares of the weights (L2 norm)

Equation 4-10. Lasso Regression cost function

$$J(\theta) = \text{MSE}(\theta) + \alpha \sum_{i=1}^n |\theta_i|$$



Performance



Regularization

That's it for me!

Over to Tom Beucler now!

3) How are ML models trained & evaluated?

4) What is physics-guided ML?



Physics-Guided ML: Add physical structure to restrict ML output to physically-plausible solutions

$$\text{Relative Humidity} \stackrel{\text{def}}{=} \frac{\text{Vapor pressure}}{\text{Saturation vapor pressure}}$$



Physics-Guided ML: Add physical structure to restrict ML output to physically-plausible solutions

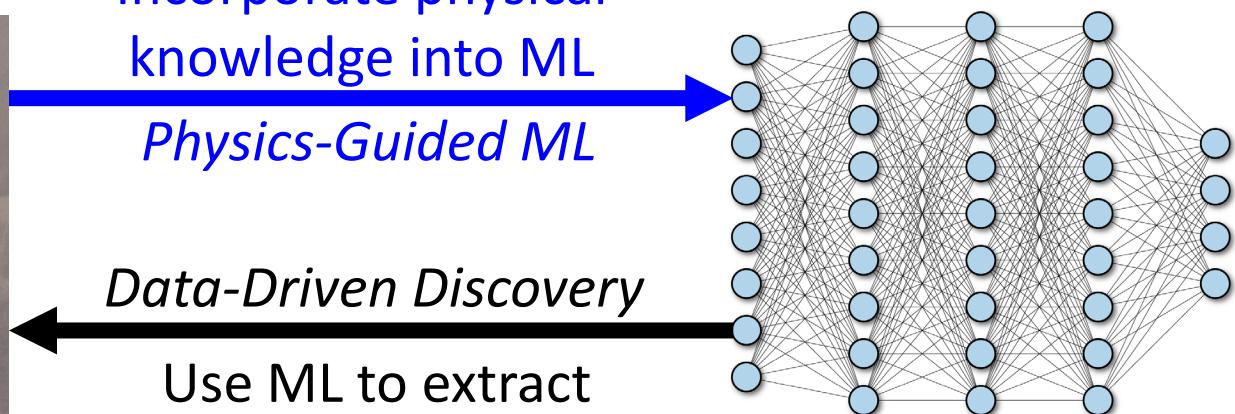
$$\frac{D\vec{v}}{Dt} + 2\vec{\Omega} \times \vec{v} = -\frac{\vec{\nabla}p}{\mathcal{C}} - \vec{\nabla}\Phi$$
$$\mathcal{M} \frac{dI}{dt} = I - B$$
$$\frac{de^*}{dT} = -\frac{\mathcal{L}_v e^*}{R_v T^2}$$

Incorporate physical knowledge into ML

Physics-Guided ML

Data-Driven Discovery

Use ML to extract physical knowledge from data



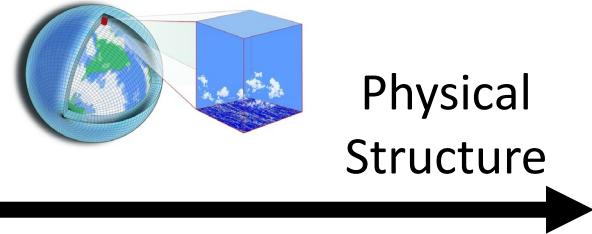
Advantages

- Physical consistency
small violations physical def./laws
- Robustness (control)
stability to small perturbations
- Improved Generalization
extrapolation to regimes/extremes
- Data efficiency
ability to learn with few samples
- Interpretability
explainable structure & behavior
- Reduced error*

Limitations

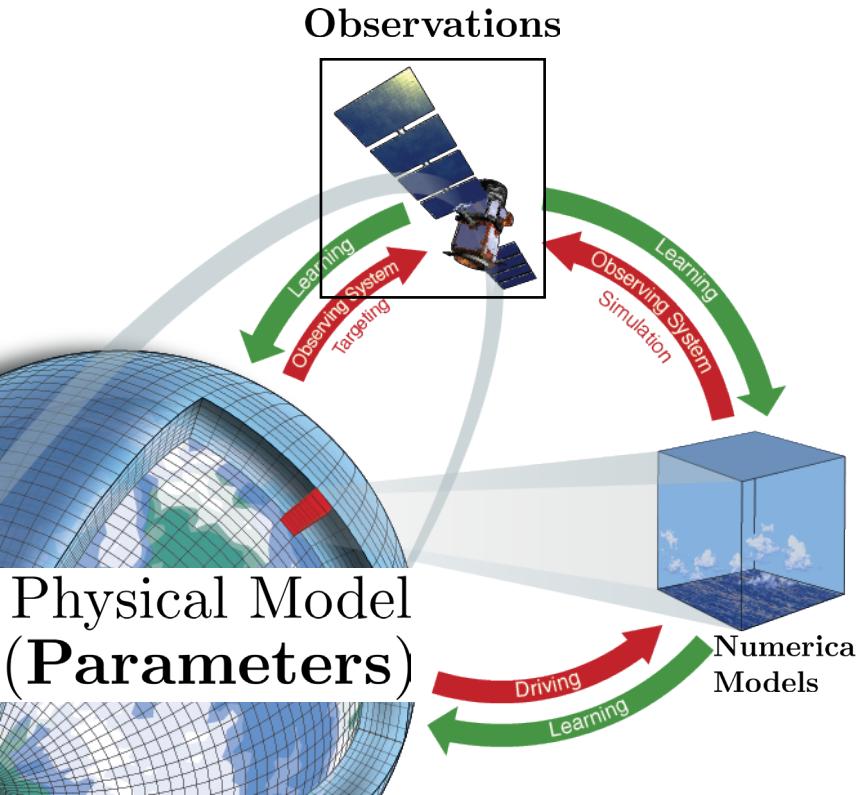
- Only a few laws hold exactly:
definitions, state eq., cons. laws
- Sources and sinks may be difficult to quantify in real data
importance of uncertainty quant.
- Trade-off between physical consistency and optimization
“no free lunch” especially when $\text{dim}(\text{Constraints}) \ll \text{effective dim.}$

Physics-Guided ML: Add physical structure to restrict ML output to physically-plausible solutions



Learn
Parameters of
Physical Model

(Most constrained) Learn Parameters of Physical Models



Parameter estimation/calibration problem

Subfield of data assimilation = Optimal state estimation given real-time data

See ECMWF resources on 4D-Var

$$p(\text{Parameters} | \text{Obs})$$

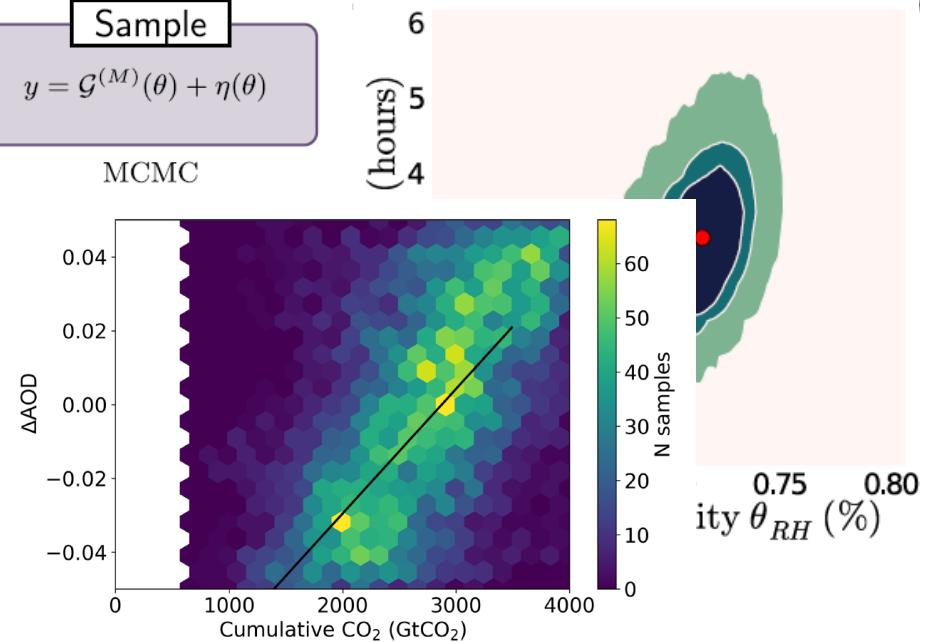
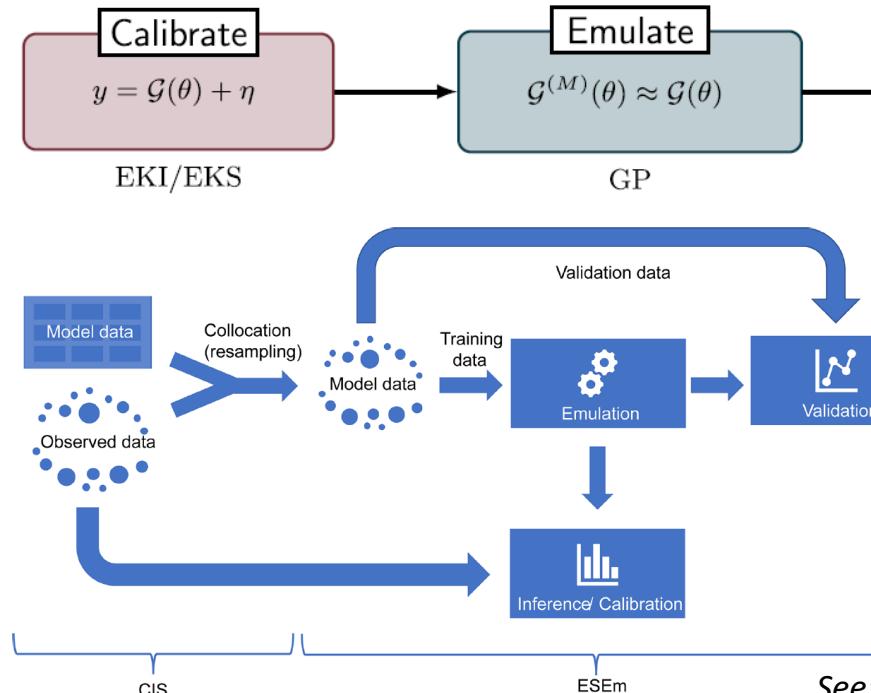
ML helps maximize the likelihood

See: Brajard et al. (2021), ECMWF Fact sheets; Image Source: CliMA

TROPICANA Physics-Guided ML Tutorial

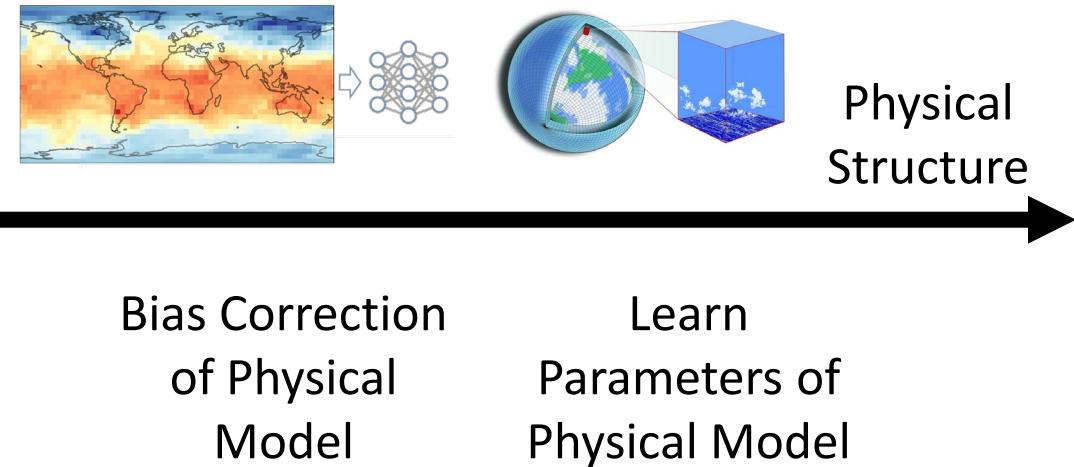
(Most constrained) Learn Parameters of Physical Models

ML-based frameworks: CES & ESEm



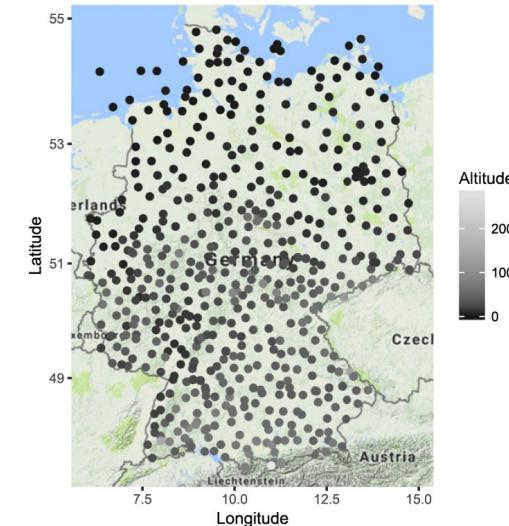
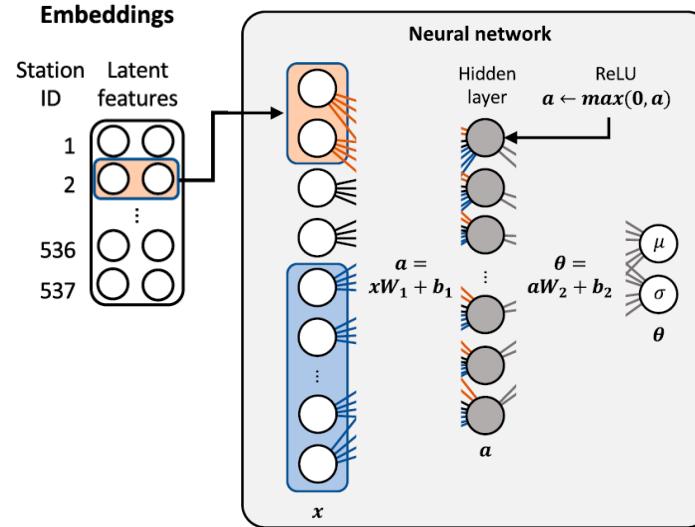
See: Cleary et al. (2021), Howland et al. (2022), Watson-Parris et al. (2021)

Physics-Guided ML: Add physical structure to restrict ML output to physically-plausible solutions



(Usually constrained) Bias-correcting physical model

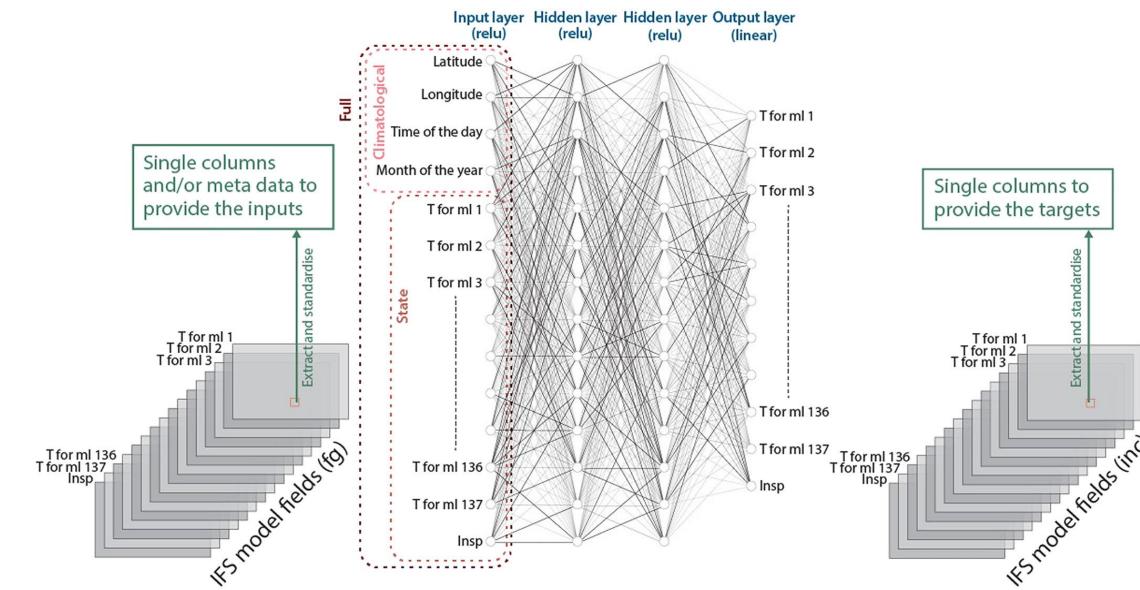
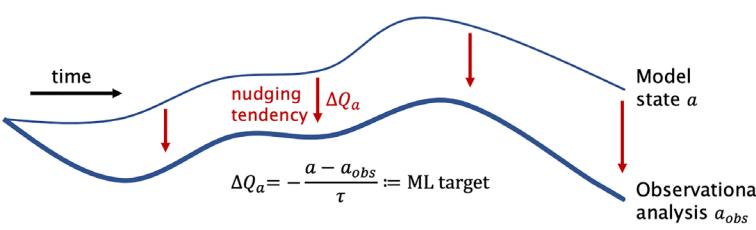
- Residual modeling: Debiased Model = (Model to correct) + **ML**
(1) Post-processing (Model Output Statistics)



See: Rasp and Lerch (2018), other ML post-processing references

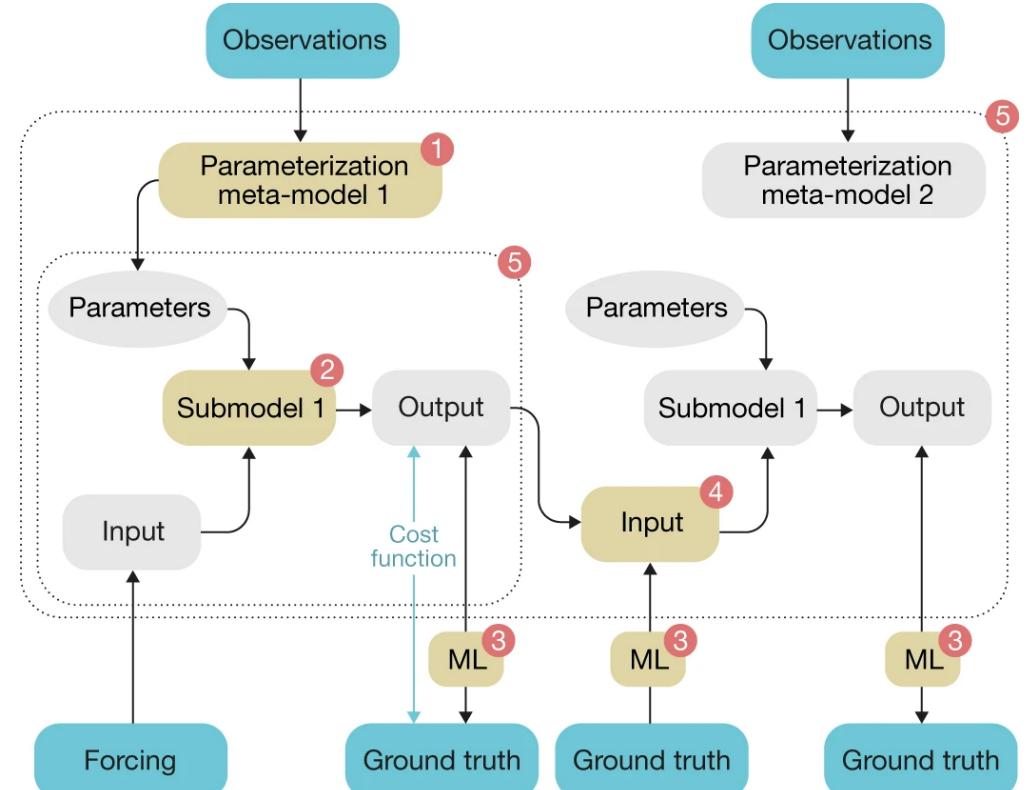
(Usually constrained) Bias-correcting physical model

- Residual modeling:
 - (1) Post-processing
 - (2) Nudging/DA increments

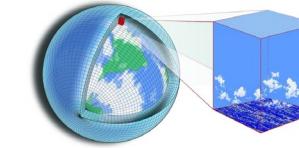
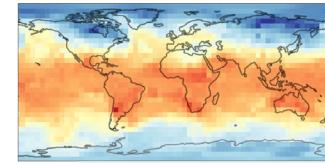
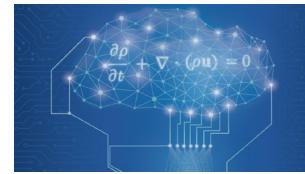


(Usually constrained) Bias-correcting physical model

- Residual modeling:
(1) Post-processing
(2) Nudging /DA increments
- Hybrid modeling: Combining physical model & ML



Physics-Guided ML: Add physical structure to restrict ML output to physically-plausible solutions



Physical
Structure

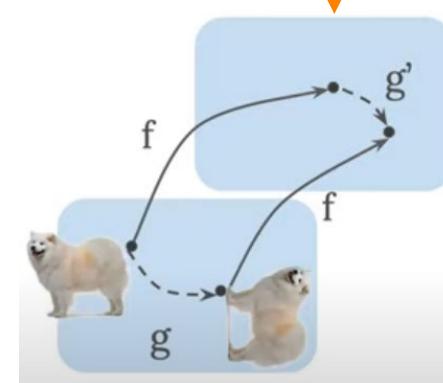
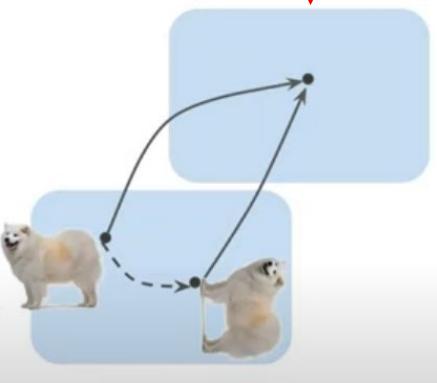
Physics-
Constrained Loss
or Architecture

Bias Correction
of Physical
Model

Learn
Parameters of
Physical Model

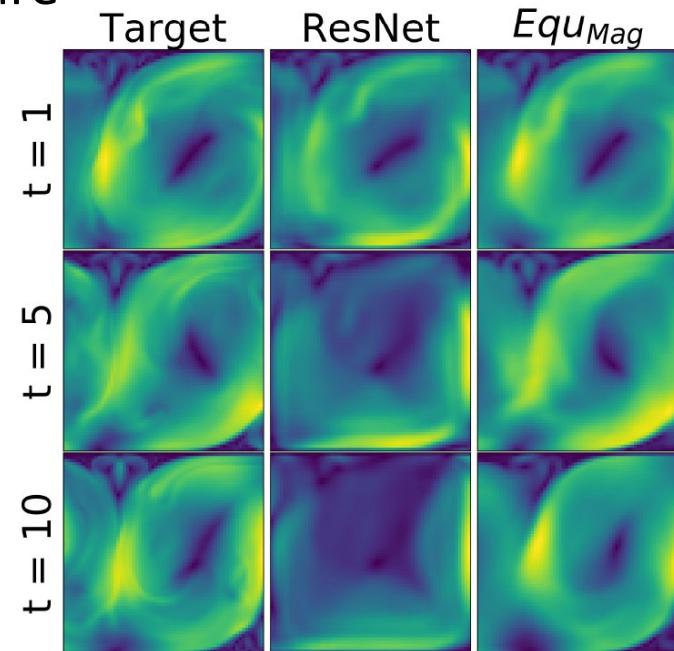
(Hard constraint) Physics-constrained architecture

- Encode **invariance** or **equivariance** in architecture



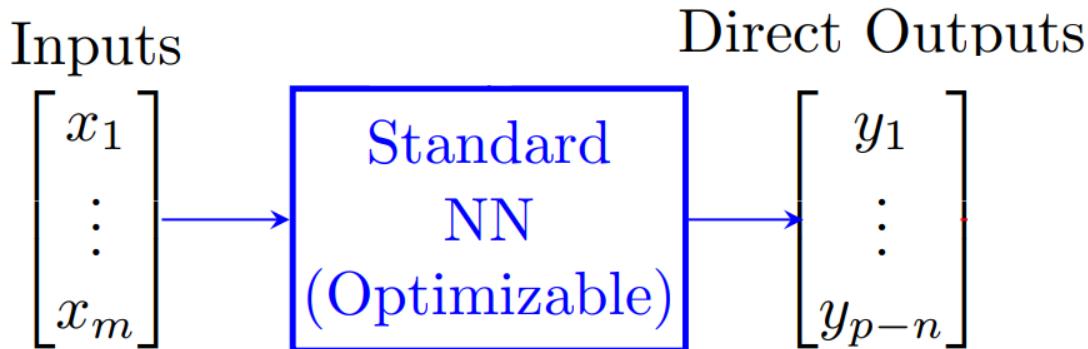
$f [g(x)] = f(x)$
Pooling layers
(large pooling size)

$f [g(x)] = g'(x)$
Convolutional
layers



(Hard constraint) Physics-constrained architecture

- Encode invariance or equivariance in architecture
- Enforce physical constraints via e.g., constraints layers



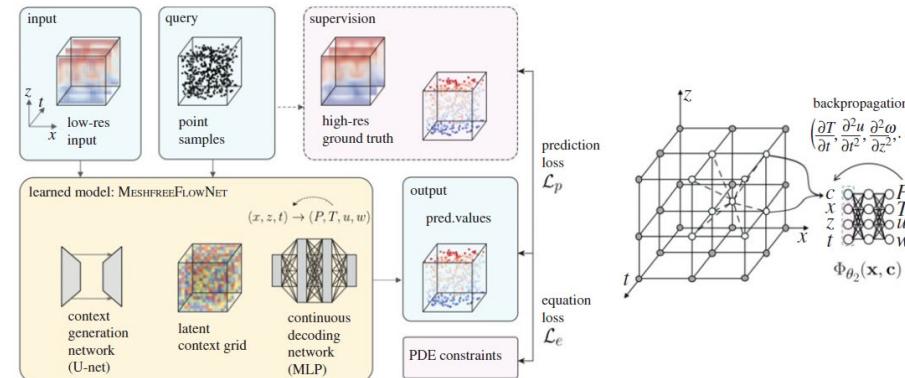
(Soft constraint) Physics-constrained loss

$$\text{Loss} = \alpha (\text{Regression loss}) + (1 - \alpha) (\text{Physics - constrained loss}) \quad \alpha \in [0, 1]$$

Physics-constrained losses:

$$\|\mathcal{P}(y_{\text{Pred}}, x)\|$$

- Residual of definitions, state equations, conservation laws



See: Jiang et al. (2020)

(Soft constraint) Physics-constrained loss

Loss = α (Regression loss) + $(1 - \alpha)$ (Physics – constrained loss) $\alpha \in [0, 1]$

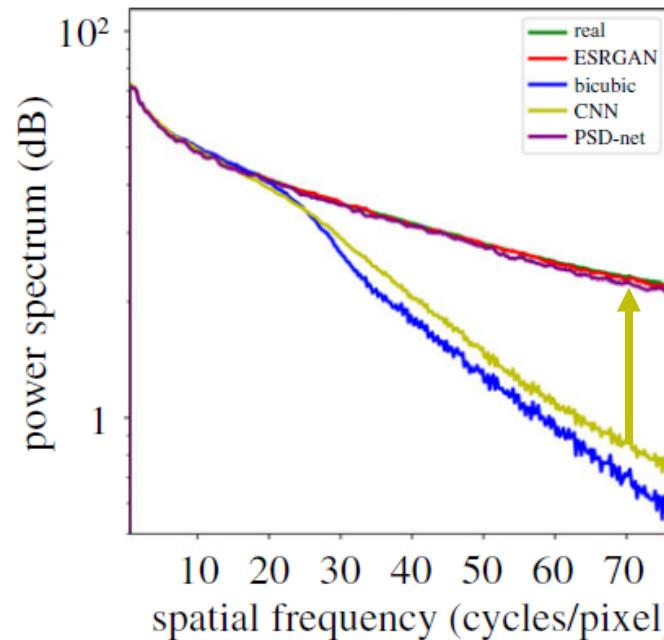
Physics-constrained losses:

$$\|\tilde{\mathcal{P}}(y_{\text{Pred}}, x)\|$$

- Residual of definitions, state equations, conservation laws
- Approximate dynamics

(Soft constraint) Physics-constrained loss

$$\text{Loss} = \alpha (\text{Regression loss}) + (1 - \alpha) (\text{Physics - constrained loss}) \quad \alpha \in [0, 1]$$



Regression loss $[\widehat{y_{\text{pred}}}(\mathbf{k}, \omega), \widehat{y_{\text{true}}}(\mathbf{k}, \omega)]$

e equations, conservation laws

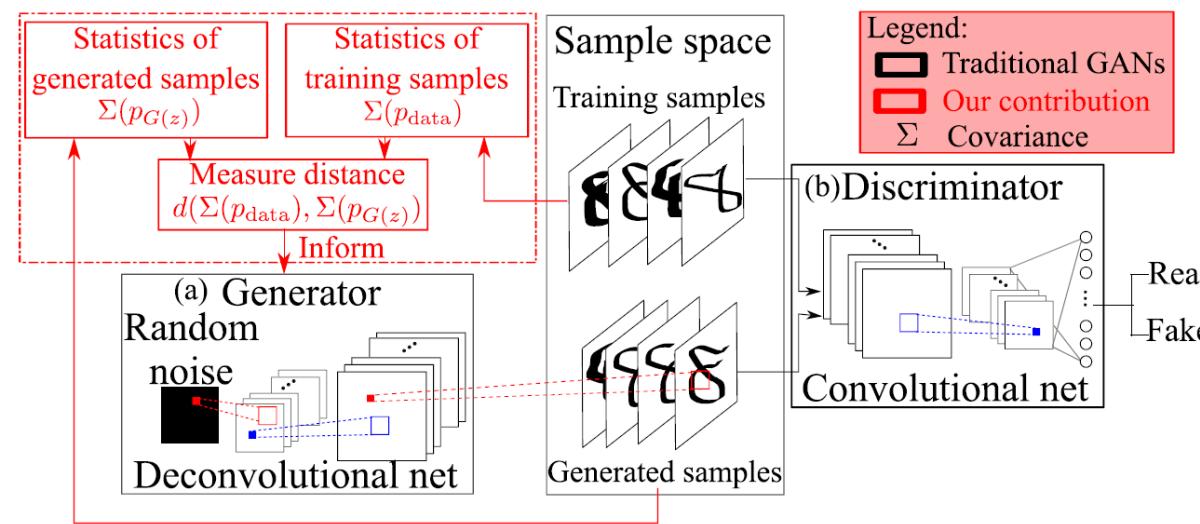
oss; (2) directly predict spectral coefficients

(Soft constraint) Physics-constrained loss

$$\text{Loss} = \alpha (\text{Regression loss}) + (1 - \alpha) (\text{Physics - constrained loss}) \quad \alpha \in [0, 1]$$

Physics-constrained losses:

$\|\text{Covariance}(y_{\text{pred}}) - \text{Covariance}(y_{\text{true}})\|$



variation laws

ict spectral coefficients
omments to truth

See: Wu et al. (2018), Mooers et al. (2020)

(Soft constraint) Physics-constrained loss

$$\text{Loss} = \alpha (\text{Regression loss}) + (1 - \alpha) (\text{Physics - constrained loss}) \quad \alpha \in [0, 1]$$

Physics-constrained losses:

$$\text{Regression loss} \left[\begin{pmatrix} \text{Mean}_{\text{pred}} \\ \text{Std}_{\text{pred}} \\ \text{Skewness}_{\text{pred}} \\ \dots \end{pmatrix}, \begin{pmatrix} \text{Mean}_{\text{true}} \\ \text{Std}_{\text{true}} \\ \text{Skewness}_{\text{true}} \\ \dots \end{pmatrix} \right]$$

- Residual of definitions, state equations, conservation laws
- Approximate dynamics
- (1) Multi-scale or spectral loss; (2) directly predict spectral coefficients
- Statistical constraints: (1) Compare statistical moments to truth
(2) Predict moments of prescribed distributions

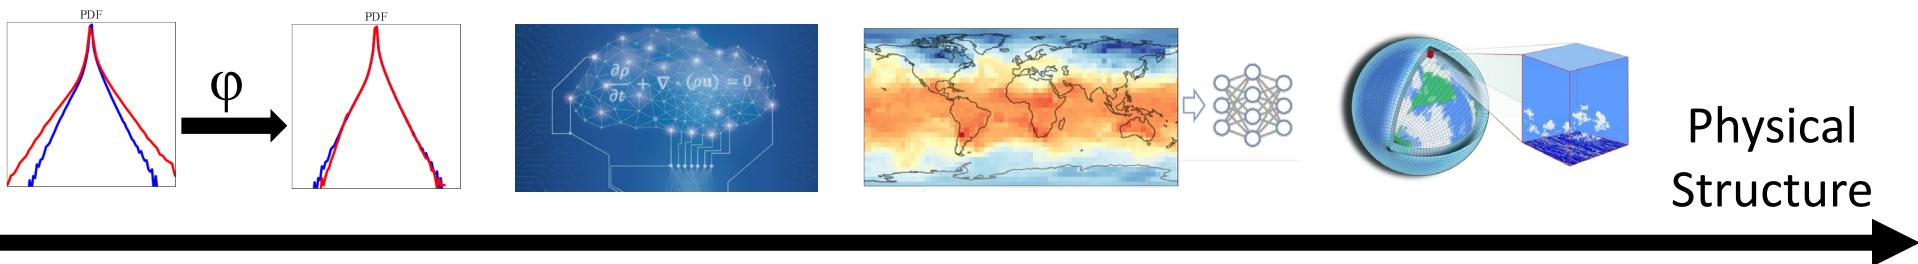
(Soft constraint) Physics-constrained loss

$$\text{Loss} = \alpha (\text{Regression loss}) + (1 - \alpha) (\text{Physics - constrained loss}) \quad \alpha \in [0, 1]$$

Physics-constrained losses: Regression loss $[\mathcal{W}(\mathbf{y}_{\text{pred}}, \mathbf{y}_{\text{true}})]$

- Residual of definitions, state equations, conservation laws
- Approximate dynamics
- (1) Multi-scale or spectral loss; (2) directly predict spectral coefficients
- Statistical constraints: (1) Compare statistical moments to truth
(2) Predict moments of prescribed distributions
- Custom weighting to emphasize certain physical regimes

Physics-Guided ML: Add physical structure to restrict ML output to physically-plausible solutions



Physics-Informed
Feature selection
/transformation

Physics-
Constrained Loss
or Architecture

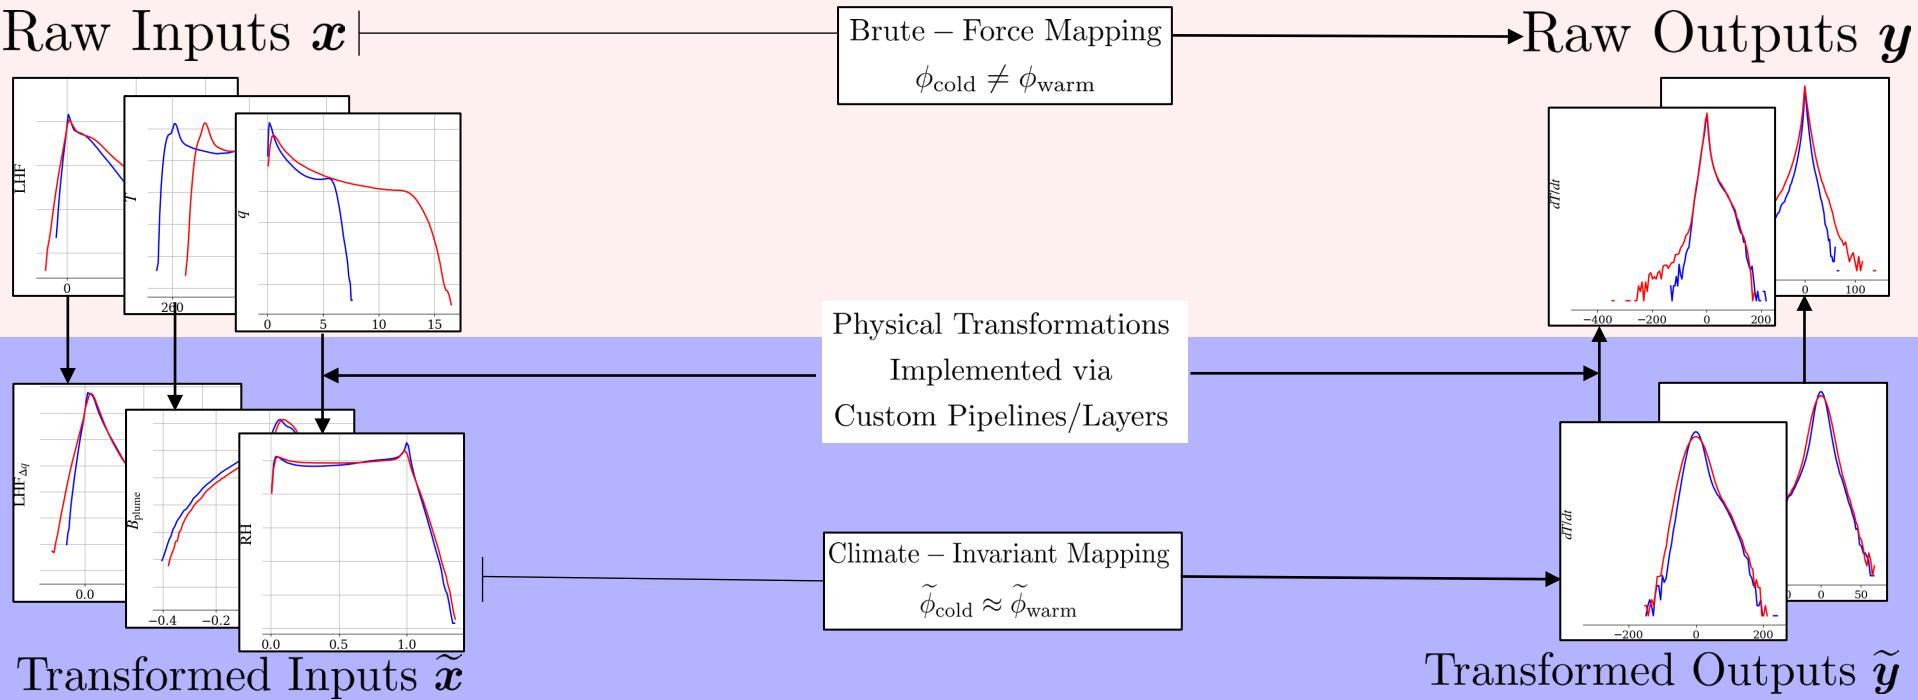
Bias Correction
of Physical
Model

Learn
Parameters of
Physical Model

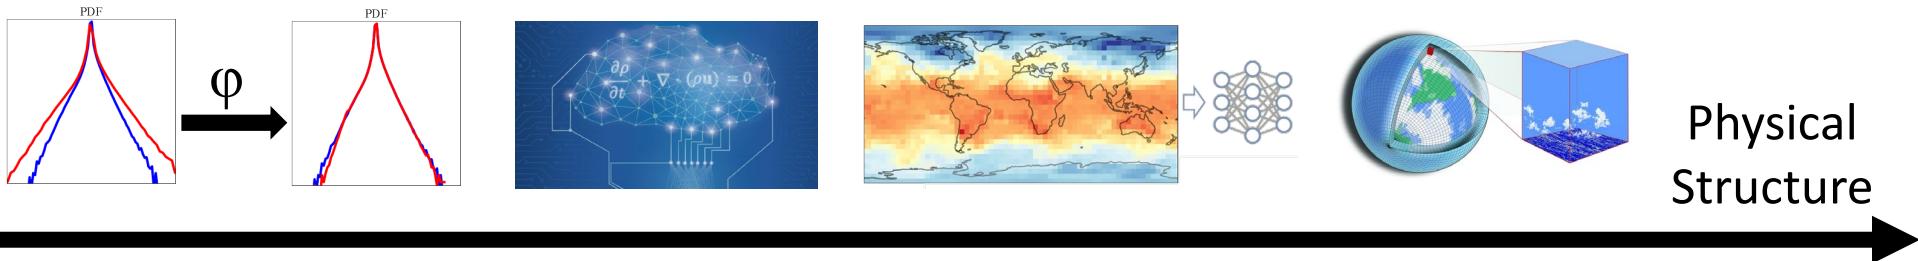
TROPICANA Physics-Guided ML Tutorial

(Guided) Physics-informed feature selection/transformation

Feature selection: Using physical knowledge, causal discovery algorithms...



Physics-Guided ML: Add physical structure to restrict ML output to physically-plausible solutions

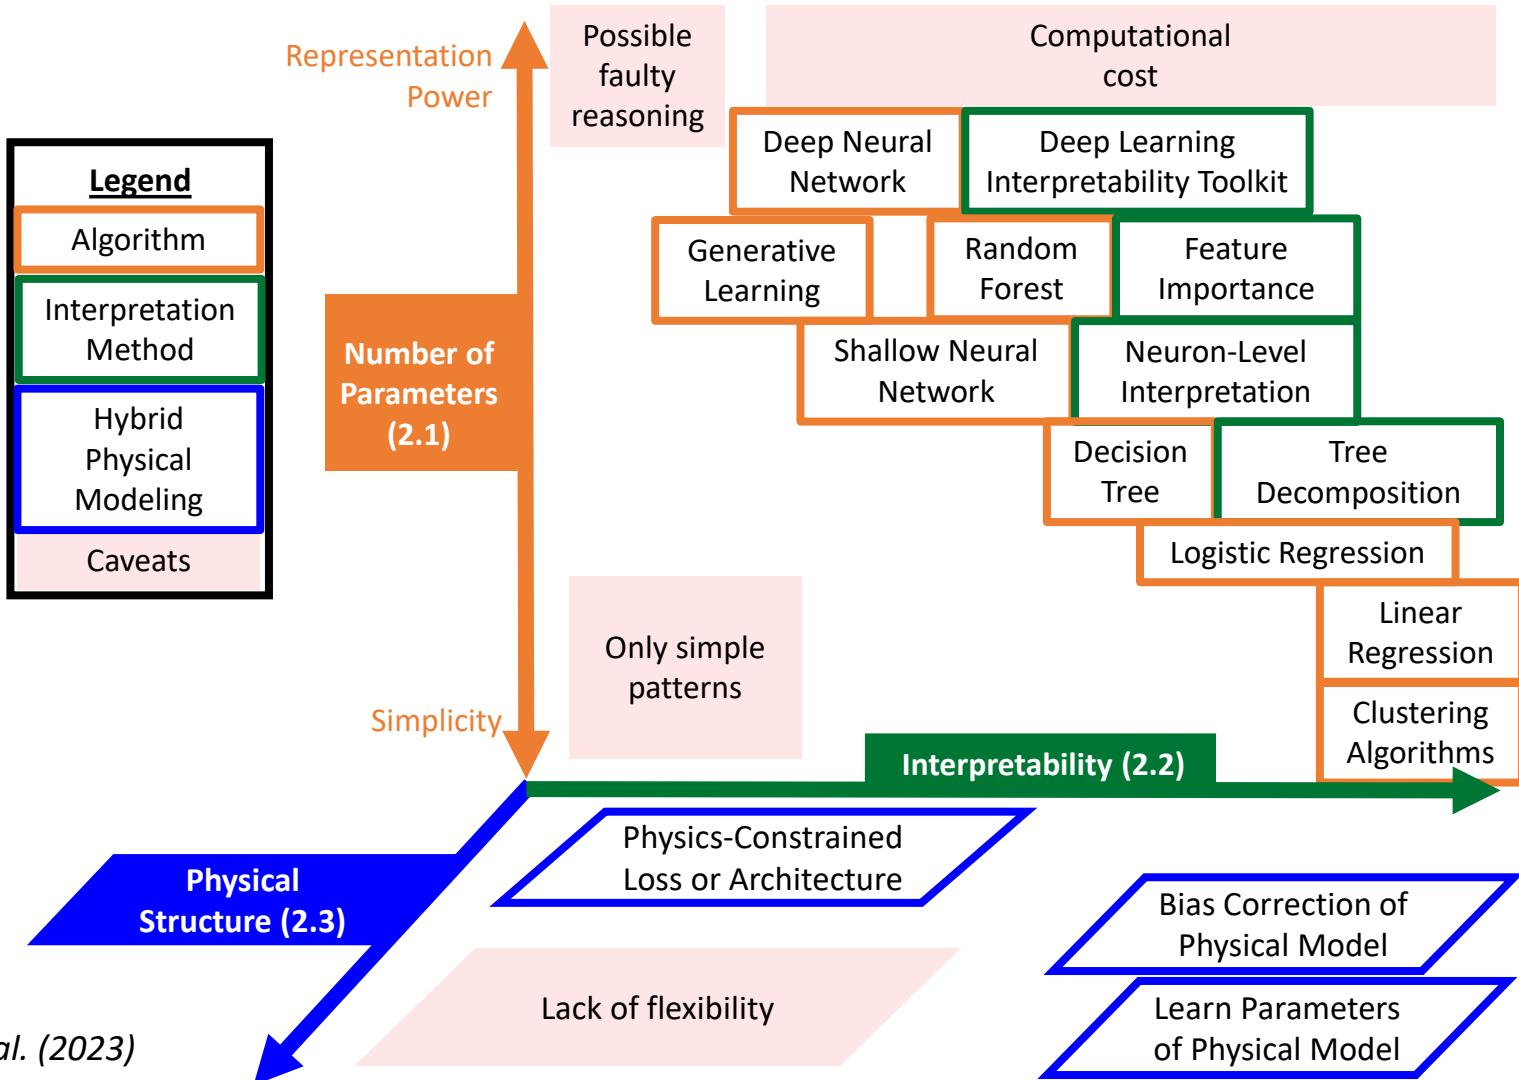


Physics-Informed
Feature selection
/transformation

Physics-
Constrained Loss
or Architecture

Bias Correction
of Physical
Model

Learn
Parameters of
Physical Model



TROPICANA GitHub Repository:

https://github.com/tbeucler/2024_TROPICANA_ML

TROPICANA GitHub Repository:

https://github.com/tbeucler/2024_TROPICANA_ML